
Conformalized Heteroscedastic Variant Effect Prediction

with Local Coverage Guarantees

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We introduce **HCCP** (Heteroscedastic Class-Conditional Conformal Prediction),
2 a post-hoc framework that pairs a learned variance head $\hat{\sigma}(x)$ with Mondrian-
3 ($y \times \hat{\sigma}$ -bin) calibration via a feature-dependent nonconformity score, attaining
4 class-conditional and bin-local coverage simultaneously. For binary classification
5 under both class imbalance and prediction-dependent noise — the regime of clinical
6 variant effect prediction (VEP), where TraitGym [Benegas et al., 2025] minority
7 prevalence sits at $\pi_{\min} \approx 0.10$ — existing conformal methods trade these two
8 objectives, a Pareto wall sharpened by the finite-sample pointwise impossibility
9 of Barber et al. [2020]. Within the equi-bin Mondrian- K family we prove an
10 $O(n^{-1/2})$ upper bound on the worst-cell coverage gap, matched within this family
11 by an $\Omega(n^{-1/2})$ lower bound with $\pi_{\min}^{-1/2}$ explicit in the constant (T5.1, T5.2; The-
12 orems 3, 4); the bound is dimension-free given $\hat{\sigma}$ is 1-D Lipschitz. On TraitGym
13 Complex ($n = 11,400$), HCCP yields disjoint-CI improvements over $\hat{\sigma}$ -Mondrian
14 [Boström and Johansson, 2020], class-Mondrian, RLCP, weighted CP, and SC-CP
15 at matched $(\hat{p}, \hat{\sigma})$ (recommended operating point $K_{\text{eval}} = 5$, per-cell minority
16 ≥ 100 ; $B = 200$ chromosome-bootstrap), while on Mendelian ($n = 3,380$) the
17 gain is restricted to a $K_{\text{eval}} = 3$ window — a regime boundary predicted by T5.1
18 (Theorem 3) when per-cell minority count falls below the heteroscedastic head’s
19 identifiability threshold. We replicate cross-platform on Open Targets matched-9
20 [Ochoa et al., 2023] (feature-pipeline-disjoint; $16.7\times$ at the recommended operat-
21 ing point with the same per-cell-minority-driven K_{eval} reversal) and cross-domain
22 on ProteinGym [Notin et al., 2023].

23 1 Introduction

24 Clinical variant interpretation drives consequential decisions [Richards et al., 2015], yet machine-
25 learned variant effect predictors (VEPs) collapse evidence into a single probability [Kircher et al.,
26 2014, Linder et al., 2025, Benegas et al., 2023, 2025] without explicit confidence. Conformal
27 prediction guarantees distribution-free coverage [Vovk et al., 2005, 2003, Romano et al., 2019, Barber
28 et al., 2021, 2023], yet Barber et al. [2020] rule out pointwise conditional coverage in finite samples.
29 VEP therefore needs a method that (i) consumes a learned heteroscedastic head $\hat{\sigma}(x)$ as in DEGU
30 [Zhou et al., 2026], (ii) stratifies calibration by both label and $\hat{\sigma}$ -bin to recover as much conditioning as
31 finite samples permit, and (iii) holds local coverage across distribution shifts. The gap binds for VEP:
32 TraitGym minority prevalence sits at $\pi_{\min} \approx 0.10$, and chrom-LOO, trait-LOO, and cross-dataset
33 shifts are intrinsic.




Figure 1: HCCP at a glance. (a) Per-cell empirical coverage on TraitGym Complex (CADD+GPN-MSA+Borzoi, $\pi_+ = 0.10$, $K = 5$): three grids over $Y \times \hat{\sigma}$ -bin showing the deviation $\text{cov} - 0.90$ (red = under-coverage, blue = over-coverage; cells annotated with signed deviation). (i) Split CP collapses on both ends of $\hat{\sigma}$ (worst cell 0.67). (ii) class-Mondrian [Sadinle et al., 2019] fixes the class budget but leaves $\hat{\sigma}$ -bin variation within each class unconstrained (worst cell 0.80). (iii) HCCP partitions on $Y \times \hat{\sigma}$ -bin and holds every cell within ± 0.04 of target. The nonconformity score $s(x, y) = |y - \hat{p}(x)| / (\hat{\sigma}(x) + \epsilon)$ with per-cell calibration restores both (T2) class- and (T3) bin-conditional coverage, compatible with Barber et al. [2020]’s pointwise impossibility. (b) Pareto plot on $(\text{Cov}|Y=1, \hat{\sigma}\text{-bin gap})$ using Tab. 9 bootstrap means ($B=200$). HCCP is the only method in the ideal corner ($\text{cov} \geq 0.85$, $\text{gap} \leq 0.20$) on both datasets; B1, B2, B3 each fail at least one axis at $\pi_+ = 0.10$.

34 The technical question is: which Mondrian partition simultaneously delivers (i) class-conditional
 35 coverage under extreme imbalance and (ii) bin-local coverage under heteroscedasticity? Single-axis
 36 answers fail individually: $\hat{\sigma}$ -Mondrian [Boström and Johansson, 2020] attains local coverage in shape
 37 but loses minority-class coverage ($\text{cov}|Y=1 = 0.62$ on TraitGym Complex), while class-Mondrian
 38 [Sadinle et al., 2019] attains class coverage but leaves an order-of-magnitude wider bin-local gap
 39 (Tab. 9). Within the equi-bin Mondrian- K family we prove the joint $(y \times \hat{\sigma}\text{-bin})$ Mondrian attains both
 40 with a finite-sample $O(\sqrt{L_F R / (\pi_{\min} n)})$ rate (T5.1, T5.2; Theorems 3, 4); $\pi_{\min}$ enters *explicitly* as
 41 a $\pi_{\min}^{-1/2}$ constant.

42 Contributions.

- 43 1. **Finite-sample $O(n^{-1/2})$ rate matched within the equi-bin Mondrian- K class with $\pi_{\min}$**
 44 **explicit in the constant** (T5.1, T5.2; Theorems 3, 4). To our knowledge this is the first finite-
 45 sample bound for class-conditional + heteroscedastic-local CP in which the minority-class
 46 probability $\pi_{\min}$ appears in the rate; the constant is dimension-free under a 1-D Lipschitz
 47 assumption on $\hat{\sigma}$. Concurrent imbalanced-CP work [Xie et al., 2025, Ding et al., 2023, Shi
 48 et al., 2024] provides marginal or class-cond rates without the $\pi_{\min}$ -local product.
- 49 2. **HCCP, a post-hoc framework with an exchangeability-robust corollary.** HCCP transfers
 50 the heteroscedastic head of DEGU from MPRA regression to binary variant classification
 51 with finite-sample local coverage, pairing a nonconformity score $s(x, y) = |y - \hat{p}(x)| / \hat{\sigma}(x)$
 52 with Mondrian- $(y \times \hat{\sigma}\text{-bin})$ calibration. T3 gives exact bin-conditional coverage at rate
 53 $1/(n_{k,b} + 1)$; T3' (via Barber et al. [2023] Thm. 2 at the cell level) supplies a uniform finite-
 54 sample lower bound when our per-chromosome KS audit detects cell-level exchangeability
 55 violations on Mendelian.
- 56 3. **TraitGym + Open Targets cross-platform validation.** At the recommended operating
 57 point ($K_{\text{eval}} = 3/5$; per-cell minority ≥ 100), HCCP is the only single-fold partition we
 58 evaluated (across split CP, $\hat{\sigma}$ -Mondrian, class-Mondrian, RLCP, weighted CP, and SC-CP at
 59 matched $(\hat{p}, \hat{\sigma})$) that holds $\text{cov}|Y=1 \geq 0.85$ and $\hat{\sigma}\text{-bin gap} \leq 0.20$ on both datasets ($B = 200$
 60 chrom-bootstrap CIs); a feature-pipeline-disjoint Open Targets matched-9 replication [Ochoa
 61 et al., 2023] reproduces the $16.7\times$ HCCP win and the K_{eval} reversal at per-cell-minority
 62 < 120 (App. G). The strong/soft asymmetry between the two TraitGym datasets (Complex:

100% paired-resample wins; Mendelian: 74.5% wins, CI crosses zero) is itself a regime boundary anticipated by the T5.1 bias-variance decomposition, not a model-class limitation.

4. **Cross-domain and synthetic validation.** ProteinGym [Notin et al., 2023] 50-assay-LOO confirms cross-domain robustness (marginal coverage 0.887, App. E); a synthetic n -sweep verifies the dimension-free $O(n^{-1/2})$ rate of T5.1 (App. F). The conformal-side gain is decoupled from base-predictor accuracy (App. C.5).

2 Related work

Conformal prediction with conditioning. Split / Mondrian CP [Vovk et al., 2005, 2003, Angelopoulos and Bates, 2022] provides marginal / group-conditional coverage; Barber et al. [2020] rule out pointwise conditional coverage. Normalized scores with a learned $\hat{\sigma}$ originate with Papadopoulos et al. [2008] and Lei et al. [2018] for regression; Boström and Johansson [2020] systematized Mondrian conformal regression with normalized scores; CQR [Romano et al., 2019] is the regression analog of our score. Recent local-CP work targets feature-space coverage at slower rates: **RLCP** [Hore and Barber, 2025] achieves $O(n^{-2/(d+2)})$ kernel-localized coverage (intractable at $d \approx 7,700$); **Self-Calibrating CP** [van der Laan, 2024] achieves $O(n^{-2/3})$ on the $\hat{\mu}$ -axis via isotonic monotonicity [Deng et al., 2021], on a different axis from ours. Tibshirani et al. [2019], Gibbs and Candès [2021], Barber et al. [2023] cover distribution-shift robustness; we use Barber 2023 Thm. 2 for our T3'. De Wolf et al. [2025] analyze conditional validity of heteroscedastic CP for regression; we extend to classification with $\pi_{\min}$ explicit. Yao et al. [2026] (concurrent) prove a dimension-free $O(n^{-1/2})$ efficiency rate for split-CP quantile regression trained via SGD; their result establishes the rate on linear quantile reg. We contribute orthogonal axes — the partition-family analysis (equi-bin Mondrian- K), explicit $\pi_{\min}^{-1/2}$ constant, and within-class tight lower bound — not a faster rate. Plassier et al. [2025] target approximate conditional validity via density-ratio reweighting under TV-distance; an alternative to discrete $\hat{\sigma}$ -bin Mondrian without finite-sample bin-conditional guarantees. Toccaceli [2026] (concurrent) identifies optimal bin count under covariate-sorted contiguous binning as an open problem; Theorem 3 resolves a special case for $(y \times \hat{\sigma}\text{-bin})$ Mondrian under class imbalance.

Class-conditional and imbalanced CP. Sadinle et al. [2019] introduced label-conditional CP (Mondrian-by- y); Shi et al. [2024], Ding et al. [2023] adapt class-conditional CP for many classes / small per-class samples; Xie et al. [2025] (concurrent) study open-set + imbalanced classification with reweighting and obtain a $O(n^{-1/2})$ marginal-coverage bound *without* $\pi_{\min}$ -local product. Bairaktari et al. [2025] (concurrent) extend group-conditional CP to overlapping/fractional groups with minimax-optimal coverage subsuming class-conditional and Mondrian as static-group special cases; their framework requires pre-specified groups and does not cover our data-dependent $(y \times \hat{\sigma}\text{-bin})$ partition or provide a matching lower bound. None of these address the heteroscedastic axis: in our H2H ablation (Tab. 9), single-axis Mondrian predictors face a fundamental tradeoff in the imbalanced regime — class-Mondrian preserves $\text{cov}_{|Y=1}$ but loses bin-local coverage (gap 0.42–0.80, bootstrap mean), while $\hat{\sigma}$ -Mondrian preserves local coverage but loses $\text{cov}_{|Y=1}$ down to 0.62 at $\pi_{\min} = 0.10$. HCCP’s joint $(y \times \hat{\sigma}\text{-bin})$ Mondrian is, to our knowledge, the first construction simultaneously satisfying both objectives in this regime, with T5 quantifying the price as $O(\sqrt{L_F R}/(\pi_{\min} n))$.

Heteroscedastic UQ and VEP benchmarks. Nix and Weigend [1994] originated Gaussian NLL for learned variance; Lakshminarayanan et al. [2017] introduced deep ensembles. **DEGU** [Zhou et al., 2026] pioneered ensemble distillation to a Gaussian-uncertainty student for massively parallel reporter assay (MPRA) regression and motivated conformal prediction as future work; we extend its $\hat{\sigma}$ machinery to (a) variant classification and (b) finite-sample local coverage with adaptive partition granularity. We do not claim a head-to-head defeat of DEGU: DEGU is a regression pipeline on MPRA with a Keras CNN ensemble, whereas HCCP targets binary classification on TraitGym, Open Targets, and ProteinGym; the “DEGU-lite” ablation in App. C reimplements the ensemble-distilled $\hat{\sigma}$ at our matched GBM backbone solely to isolate the choice of variance source, not to replicate DEGU’s published pipeline. **TraitGym** [Benegas et al., 2025] standardizes VEP evaluation; features include CADD [Kircher et al., 2014], Borzoi [Linder et al., 2025], GPN-MSA [Benegas et al., 2023]; AlphaGenome [Avsec et al., 2026] is a recent foundation model used in our App. C ablation. CP applications in genomic medicine [Papangelou et al., 2025, Wang et al., 2025, Adegoke et al., 2025] target drug response or polygenic risk scores in a regression / clinical-screening setting; HCCP

is, to our knowledge, the first CP framework with formal class-conditional and bin-local coverage guarantees for binary variant pathogenicity.

Distinguishing claim vs. concurrent work. The four most adjacent 2025–2026 results each cover a strict subset: Yao et al. [2026] dim-free $O(n^{-1/2})$ for split-CP quantile reg (no Mondrian / $\pi_{\min}$ / lower bound); Plassier et al. [2025] TV-reweighted approximate conditional validity (no finite-sample bin-cond bound); Bairaktari et al. [2025] minimax overlapping groups (pre-specified partitions only, no matching lower bound); Xie et al. [2025] marginal rate without $\pi_{\min}$ -local product. *None jointly delivers $(y \times \hat{\sigma})$ data-dependent partition + $\pi_{\min}$ -explicit constant + matching within-class lower bound; HCCP is the unique Tab. 8 row satisfying all six axes.*

3 Problem formulation

3.1 Setup and assumptions

Let $\mathcal{D} = \{(X_i, Y_i, C_i)\}_{i=1}^n$ with $X_i \in \mathcal{X}$, $Y_i \in \{0, 1\}$, and chromosome label $C_i \in \mathcal{C}$. We use the chrom-LOO protocol of Benegas et al. [2025]: a test chromosome $c^* \in \mathcal{C}$ is held out and the calibration fold $\mathcal{D}^{(-c^*)} := \{i : C_i \neq c^*\}$ is used for both the aggregator $\hat{p}$ and the heteroscedastic head $\hat{\sigma}$. Test point $(X, Y) \sim P_{c^*}$.

Assumption 1 (A1 chrom-group exchangeability). *Variants from the same chromosome are exchangeable: $\{(X_i, Y_i) : C_i = c\} \stackrel{d}{=} \{(X_{\pi(i)}, Y_{\pi(i)}) : C_i = c\}$ for any permutation π within chromosome c .*

Assumption 2 (A1' chrom-wise i.i.d. (strengthening of A1)). $\{(X_i, Y_i) : C_i = c\} \stackrel{i.i.d.}{\sim} P_c$. T3 and T5 use A1' for the cleaner sampling argument; T3' uses only A1 for robustness.

Assumption 3 (A2 score stationarity across chromosomes). *A2-marginal: score law is chrom-invariant. A2-class: invariant conditional on $Y = k$. A2-cell: invariant conditional on $Y = k$ and $\hat{\sigma}$ -bin. Hierarchy: A2-cell $\Rightarrow$ A2-class $\Rightarrow$ A2-marginal.*

Empirically, the cell-level KS audit (App. B, Tab. 4) finds approximate stationarity on Complex (10.4% rejection across 125 tests, near the 5% null) and pronounced cell-level violations on Mendelian (46.2% rejection across 52 tests, max KS 0.434); T3' is the operational certificate on Mendelian (§5.1), T3 exact suffices on Complex.

Assumption 4 (A-SL Score-Lipschitz regularity). *For each class k , $\sup_s |F_{k,\sigma_1}(s) - F_{k,\sigma_2}(s)| \leq L_F |\sigma_1 - \sigma_2|$ for all $\sigma_1, \sigma_2 \in [\hat{\sigma}_{\min}, \hat{\sigma}_{\max}]$.*

Assumption 4 is used only by T5 (§5.2) to derive the optimal bin count K^* . Empirical audit of Assumption 4 via per-pair Lipschitz ratios on $\hat{\sigma}$ -bin pairs (App. B, Fig. 6) gives noise-adjusted upper envelope $L_F^{(95\%)} = 13.8$ (Mendelian) and 8.0 (Complex); LCLS estimator [Huang et al., 2024] gives a tighter typical-slope estimate $\hat{L}_F = 2.97 / 4.51$.

3.2 Score and prediction set

Aggregator $\hat{p}$ (HistGradientBoosting depth 2, class-balanced) and **heteroscedastic head** $\hat{\sigma}$ (same backbone, Gaussian NLL on chrom-LOO OOF residuals $r_i = |Y_i - \hat{p}^{(-C_i)}(X_i)|$). **Nonconformity score** $s(x, y) := |y - \hat{p}(x)| / (\hat{\sigma}(x) + \epsilon)$, $\epsilon = 10^{-6}$. **Mondrian prediction set:** partition $\mathcal{D}^{(-c^*)}$ into $2K$ cells via $(Y_i, \text{bin}_K(\hat{\sigma}(X_i)))$ where bin_K uses calibration-fold $\hat{\sigma}$ -quantiles. Per-cell threshold:

$$\hat{q}_{k,b} := \text{Quantile}(\{s(X_i, Y_i) : Y_i = k, \text{bin}_K(\hat{\sigma}(X_i)) = b\}, \lceil (n_{k,b} + 1)(1 - \alpha) \rceil / n_{k,b}). \quad (1)$$

The Mondrian prediction set:

$$\mathcal{C}_\alpha(x) := \{k \in \{0, 1\} : |k - \hat{p}(x)| \leq \hat{q}_{k, \text{bin}_K(\hat{\sigma}(x))} \cdot (\hat{\sigma}(x) + \epsilon)\}. \quad (2)$$

For cells with $n_{k,b} < \lceil 1/\alpha \rceil$ we fall back to the pooled class-conditional threshold $\hat{q}_k$ (T2).

4 Method: HCCP

Pipeline. HCCP is a three-stage post-hoc framework on top of any probabilistic classifier: (1) aggregator $\hat{p}$, (2) heteroscedastic head $\hat{\sigma}$, (3) Mondrian- $(y \times \hat{\sigma}$ -bin) conformal calibration. Each

Algorithm 1 HCCP — Heteroscedastic Class-Conditional Conformal Prediction

Require: Data $\mathcal{D}$, test chrom c^* , miscoverage α , bin count K

```

1:  $\mathcal{D}_{\text{tr}} \leftarrow \{i : C_i \neq c^*\}$ ; fit  $\hat{p}$  on  $\mathcal{D}_{\text{tr}}$  with chrom-LOO OOF predictions  $\{\hat{p}^{(-C_i)}(X_i)\}$ 
2: Fit  $\hat{\sigma}$  on  $\mathcal{D}_{\text{tr}}$  with targets  $\{|Y_i - \hat{p}^{(-C_i)}(X_i)|\}$  under Gaussian NLL
3: Compute  $\hat{\sigma}$ -bin edges via calibration-fold quantiles; calibration scores  $s_i = |Y_i - \hat{p}(X_i)|/(\hat{\sigma}(X_i) + \epsilon)$ 
4: Partition  $\mathcal{D}_{\text{tr}}$  by  $(Y_i, \text{bin}_K(\hat{\sigma}(X_i)))$ ; compute  $\hat{q}_{k,b}$  via Eq. (1)
5: for cells with  $n_{k,b} < \lceil 1/\alpha \rceil$  do
6:    $\hat{q}_{k,b} \leftarrow \hat{q}_k$  ▷ T2 pooled class-conditional fallback
7: end for
8: for  $x \in \mathcal{D}_{\text{te}}$  do
9:    $\mathcal{C}_\alpha(x) \leftarrow \{k \in \{0, 1\} : |k - \hat{p}(x)| \leq \hat{q}_{k,b(x)} \cdot \hat{\sigma}(x)\}$ 
10: end for
11: return  $\{(x, \mathcal{C}_\alpha(x))\}$ 

```

stage is decoupled, so HCCP accepts any $(\hat{p}, \hat{\sigma})$ pair as a drop-in. Algorithm 1 states the full pipeline; per-cell small-cell fallback to the pooled T2 threshold triggers when $n_{k,b} < \lceil 1/\alpha \rceil$.

Why the joint $(y \times \hat{\sigma}$ -bin) partition. Pure class-conditional calibration (Mondrian on y alone) [Vovk et al., 2003] fixes the imbalance pathology of split CP but leaves within-class local coverage unaddressed: our ablation reaches $\text{cov}_{|Y=1} = 0.896$ on Complex but the hardest $\hat{\sigma}$ -decile has coverage 0.454. Stratifying by $\hat{\sigma}$ -bin alone re-introduces minority-class pathology because $\hat{\sigma}$ -distribution is dominated by negatives. Among Mondrian partitions over $\{y, b(x)\}$, the product partition $(y \times b)$ is the unique refinement that admits both T2 (which requires conditioning on y) and T3 (which requires conditioning on b); the two strict coarsenings each violate one of the two coverage targets.

K selection: nested chrom-LOO. We select $\hat{K}(c_{\text{outer}})$ per outer fold via *proper nested chrom-LOO cross-validation*, ensuring $\hat{K}$ is a deterministic function of the outer calibration set alone (CP exchangeability-preserving). For each outer fold, run inner chrom-LOO and pick $\hat{K}(c_{\text{outer}}) = \text{argmin}_K \bar{G}_{\text{inner}}(K)$ where the inner-fold worst-cell gap is measured on a fixed- K_{eval} sigma-bin partition (not on each candidate K itself, which would artificially favour smaller K via fewer cells). K_{eval} is set per dataset to keep ≥ 100 minority-class samples per cell: $K_{\text{eval}} = 3$ on Mendelian, $K_{\text{eval}} = 5$ on Complex. Modal $\hat{K}(c_{\text{outer}})$ is 5 on Complex and 5–8 on Mendelian (App. C).

5 Theory

This section states HCCP’s four main-text theorems. Auxiliary lemmas (T1/T2 corollaries; T1’/T2’/T4 robust counterparts via Barber et al. [2023] Thm. 2; T3-loc feature-ball extension; T3.b $\hat{\sigma}$ -perturbation robustness) are stated and proved in App. A.

5.1 Bin-conditional coverage (T3 / T3’)

Theorem 1 (T3 bin-conditional coverage, exact). *Under Assumption 2 and A2-cell, for each cell with $n_{k,b} \geq \lceil 1/\alpha \rceil$ and continuous F_{kb} , $1 - \alpha \leq P(Y \in \mathcal{C}_\alpha(X) \mid Y = k, b(X) = b) \leq 1 - \alpha + 1/(n_{k,b} + 1)$.*

Positioning vs. Barber et al. [2020]. We do not claim pointwise conditional coverage $P(Y \in \mathcal{C}_\alpha(X) \mid X = x) \geq 1 - \alpha$, ruled out in finite samples without structural assumptions. T3 conditions on the (y, b) -partition event; the $\hat{\sigma}$ -bins are the coarsest practically useful strengthening of standard class-conditional CP that still carries finite-sample bounds.

Theorem 2 (T3’ bin-conditional coverage, robust). *Let $\bar{\delta}_{\text{TV}}^{(k,b)}$ be the average TV distance between calibration-fold and test-fold conditional score distributions in cell (k, b) . Under Assumption 2 alone (without A2-cell), for each cell with $n_{k,b} \geq \lceil 1/\alpha \rceil$, $P(Y \in \mathcal{C}_\alpha(X) \mid Y = k, b(X) = b) \geq (1 - \alpha) - \bar{\delta}_{\text{TV}}^{(k,b)}$.*

191 T3' is the operative *certificate* on Mendelian rather than a sharp estimate: empirical $\bar{\delta}_{TV}^{(k,b)} \leq 0.31$
 192 on the worst cell yields a coverage lower bound of 0.59, which is loose by construction (uniform
 193 over all distributions satisfying Assumption 2 alone), while the observed per-chromosome coverage
 194 gap remains ≤ 0.100 throughout (chr6 at 0.800; full per-chrom breakdown in App. C.7). T3'
 195 supplies the only uniform finite-sample lower bound on coverage when T3's exactness assumption
 196 (A2-cell) is empirically rejected. Our per-chromosome KS audit (App. B) shows the rejection rate is
 197 approximately K -invariant across $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$ (40–46% on Mendelian), so T3' supplies
 198 the certificate at every partition granularity.

199 5.2 Optimal bin count and within-class tight rate (T5.1, T5.2)

200 This is the principal theoretical contribution. T5.1 derives the optimal K by formalising the bias-
 201 variance tradeoff and T5.2 proves a finite-sample lower bound that is *tight within the equi-bin*
 202 *Mondrian- K family* (i.e., a within-class tightness statement, not a procedure-class minimax). The pair
 203 establishes that the $O(n^{-1/2})$ rate is tight on this class, dimension-free, and depends explicitly on
 204 $\pi_{\min}$ via the $\pi_{\min}^{-1/2}$ constant; concurrent split-CP rate analyses (e.g., Yao et al. [2026], also $O(n^{-1/2})$)
 205 dim-free, no Mondrian / $\pi_{\min}$ / lower bound) and other axes (e.g., self-calibrating CP [van der Laan,
 206 2024] on the $\hat{\mu}$ -axis with $O(n^{-2/3})$) are not subsumed and do not subsume ours.

207 Under Assumption 4, the worst-cell coverage gap decomposes as

$$G(K) := \max_{k,b} |\text{cov}_{k,b} - (1 - \alpha)| \leq \underbrace{\frac{L_F R}{K}}_{\text{stationarity bias}} + \underbrace{\frac{K}{\pi_{\min} n}}_{\text{finite-sample variance}}, \quad (3)$$

208 where $R = \hat{\sigma}_{\max} - \hat{\sigma}_{\min}$ and $\pi_{\min} = \min(\pi, 1 - \pi)$.

209 **Theorem 3** (T5.1 oracle bin count). *Under Assumptions 2 and 4 with equiprobable $\hat{\sigma}$ -bins, $K^* =$*
 210 *$\lfloor \sqrt{L_F R \pi_{\min} n} \rfloor$ minimises Eq. (3), with $G(K^*) \leq 2\sqrt{L_F R / (\pi_{\min} n)}$.*

211 **Theorem 4** (T5.2 within-class tight finite-sample lower bound on the equi-bin Mondrian- K class).
 212 *Restrict attention to conformal predictors using equiprobable $\hat{\sigma}$ -bin Mondrian partitions on the*
 213 *calibration fold (the class HCCP belongs to). There exists a distribution satisfying Assumptions 2 and*
 214 *4 such that, for any K within this class, $\min_K G(K) = \Omega(\sqrt{L_F R / (\pi_{\min} n)})$. Hence the $O(n^{-1/2})$*
 215 *rate of T5.1 is tight on this class, and the $\pi_{\min}^{-1/2}$ constant is not improvable without strengthening*
 216 *assumptions or moving outside the class. This is a within-class tightness statement: it does not rule*
 217 *out lower bounds in richer procedure classes (data-dependent partitions, weighted CP, RLCP, SC-CP);*
 218 *it certifies that the equi-bin Mondrian- K family — within which HCCP is instantiated — admits no*
 219 *asymptotically faster K -selection.*

220 **Why richer procedure classes can break this.** The $\Omega(n^{-1/2})$ ceiling is class-restrictive, not
 221 procedure-class minimax: SC-CP [van der Laan, 2024] attains $O(n^{-2/3})$ on the $\hat{\mu}$ -axis via isotonic
 222 monotonicity (a structural shape constraint on $\hat{\mu}$, not a partition refinement); data-dependent partitions
 223 [Plassier et al., 2025] and reweighting [Tibshirani et al., 2019] occupy further classes not covered by
 224 T5.2. The matching lower bound certifies only that within equi-bin Mondrian- K no asymptotically
 225 faster K -selection is obtainable without strengthening Assumption 4.

226 **Proof sketch.** For T5.1: $g'(K) = -L_F R / K^2 + 1 / (\pi_{\min} n) = 0$ gives $K^* = \sqrt{L_F R \pi_{\min} n}$. For
 227 T5.2: a linear-shift family $F_{k,\sigma}(s) = \Phi(s - L_F \sigma)$ matches both bias and variance lower bounds.
 228 Full proofs in App. A.11.

229 **Practical $\hat{K}$ via nested chrom-LOO.** T5.1 is descriptive of the asymptotic rate, not prescriptive of
 230 K at finite n . The plug-in LCLS $\hat{L}_F$ [Huang et al., 2024] gives asymptotic $K^* = 22$ (Mendelian)
 231 / 50 (Complex), but the per-cell minority floor $n_{k,b} \geq \lceil 1/\alpha \rceil$ binds well before K^* at TraitGym
 232 scale ($n \leq 11,400$); modal nested-CV $\hat{K}(c_{\text{outer}}) = 5$ Complex / $\{5, 8\}$ Mendelian (3–10 $\times$ smaller),
 233 reflecting the finite-sample fallback regime (App. C). The synthetic n -sweep (App. F) verifies
 234 the rate: log-log slope -0.44 ± 0.08 across feature dimensions $d \in \{8, 32, 128\}$ and $\pi_{\min} \in$
 235 $\{0.1, 0.5\}$ (per-curve std over 6 fits; T5.1 prediction -0.50 within 1 SD), confirming dimension-free.
 236 Practical recommendation: use nested chrom-LOO CV with K -grid $\{2, 3, 5, 8, 10\}$ for finite-sample
 237 K -selection; T5.1's K^* guides only the asymptotic regime.

238 5.3 Summary of guarantees

Table 1: HCCP coverage guarantees at a glance. Auxiliary lemmas (T1, T2, T1', T2', T3-loc, T3.b, T4) live in App. A.

Result	Assumptions	Guarantee
T3 (Thm. 1)	A1' + A2-cell	$\text{cov}_{k,b} \in [1 - \alpha, 1 - \alpha + 1/(n_{k,b} + 1)]$
T3' (Thm. 2)	A1' only	$\text{cov}_{k,b} \geq (1 - \alpha) - \bar{\delta}_{\text{TV}}^{(k,b)}$
T5.1 (Thm. 3)	A1' + A-SL	$G(K^*) \leq 2\sqrt{L_F R / (\pi_{\min} n)}, K^* = O(\sqrt{n})$
T5.2 (Thm. 4)	equi-bin Mondrian-K	$\min_K G(K) = \Omega(\sqrt{L_F R / (\pi_{\min} n)})$

239 6 Experiments

240 **Setup.** TraitGym Mendelian matched_9 ($n = 3,380, 113$ OMIM) and Complex matched_9 ($n =$
241 $11,400, 29$ traits); chrom-LOO with test chroms $\{17-22, X\}$ (TraitGym convention). Features:
242 CADD [Kircher et al., 2014], Borzoi [Linder et al., 2025], GPN-MSA [Benegas et al., 2023]. AUPRC
243 chrom-weighted, $\alpha = 0.10$. HCCP partition $K = \hat{K}(c_{\text{outer}})$ via nested chrom-LOO; reporting metric
244 on fixed K_{eval} (Mendelian $K_{\text{eval}} = 3$, Complex $K_{\text{eval}} = 5$, set per dataset to keep ≥ 100 minority-
245 class samples per cell). **Coverage and gap statistics throughout the H2H comparison are $B = 200$**
246 **chromosome-level bootstrap mean $\pm$ std for HCCP and all matched- $(\hat{p}, \hat{\sigma})$ baselines (Tab. 9,**
247 **Tab. 10).** The aggregator backbone comparison in Tab. 2 reports single-seed point estimates for the
248 non-HCCP backbones (different aggregator pipelines); HCCP-side bootstrap CIs are unchanged from
249 Tab. 9 (full setup in App. B).

250 6.1 Main results

Table 2: TraitGym chrom-LOO. HCCP at $K = 5$ for direct comparability with the split-CP and class-conditional baselines (the recommended Mondrian operating point $K_{\text{eval}} = 3/5$ for Mendelian / Complex is reported in Tab. 9). AUPRC chrom-weighted; $\hat{\sigma}$ -bin gap here is the **max-over- $\hat{\sigma}$ -bins marginal coverage deviation from 0.90** (averaged across class within each bin); the cell-level worst- (k, b) gap that appears in T5's bound is reported in Tab. 9.

Dataset	Aggregator	Calibration	AUPRC $\uparrow$	Cov $ Y=1$	$\hat{\sigma}$ -bin gap $\downarrow$	Singleton
Mendelian	LogReg [Benegas et al., 2025]	Split CP	0.752	0.89	0.54	91%
	GBM (ours)	Class-cond. (T2 only)	0.900	0.90	0.46	92%
	GBM	DEGU-lite + split	0.901	0.88	0.43	85%
	GBM	HCCP (ours)	0.900	0.894 ± 0.015	0.020 ± 0.012	0.776 ± 0.040
Complex	LogReg	Split CP	0.332	0.88	0.43	23%
	GBM	Class-cond. (T2 only)	0.353	0.90	0.51	31%
	GBM	DEGU-lite + split	0.353	0.86	0.43	20%
	GBM	HCCP (ours)	0.353	0.904 ± 0.004	0.004 ± 0.002	0.228 ± 0.019

251 The GBM aggregator alone reaches Mendelian AUPRC 0.900, a +14.8 pp absolute improvement
252 over the published LogReg baseline of Benegas et al. [2025] — the strongest TraitGym Mendelian
253 AUPRC we are aware of. Class-conditional calibration fixes the minority-class pathology of split CP
254 ($\text{cov}_{|Y=1} 0.54 \rightarrow 0.90$ on Complex) but leaves the local gap at 0.46–0.51. HCCP's joint ($y \times \hat{\sigma}$ -bin)
255 Mondrian attains a marginal $\hat{\sigma}$ -bin gap within one standard error of the finite-sample $1/(n_{k,b} + 1)$
256 floor (Tab. 2; bin-local ratios in Tab. 9). The corresponding cell-level worst- (k, b) gap (the object
257 that appears in T5's bound) is larger by construction: 0.173 ± 0.192 Mendelian and 0.060 ± 0.031
258 Complex at the recommended K_{eval} (Tab. 9); the two metrics differ by $\sim 8\times$ on Mendelian because
259 minority-cell undercoverage drives the cell-level worst, which the marginal-bin metric averages out.

260 6.2 Head-to-head against contemporary CP

261 We compare HCCP (B4) against six contemporary CP variants at matched $(\hat{p}, \hat{\sigma})$: split (B1), $\hat{\sigma}$ -
262 Mondrian (B2, [Boström and Johansson, 2020]), class-Mondrian (B3, [Sadinle et al., 2019]), RLCP

[Hore and Barber, 2025], weighted CP [Tibshirani et al., 2019], and SC-CP [van der Laan, 2024]. We adapt RLCP / SC-CP to the matched $(\hat{p}, \hat{\sigma})$ inputs HCCP uses; this isolates the partition choice from the $(\hat{p}, \hat{\sigma})$ inputs themselves and is not a comparison against these methods’ native localization axes (RLCP: feature-space kernel; SC-CP: $\hat{\mu}$ -axis isotonic; cf. Tab. 8). **At the recommended operating point ($K_{\text{eval}} = 3$ Mendelian, 5 Complex; per-cell minority ≥ 100), HCCP is the only single-fold partition we evaluated that holds $\text{cov}_{|Y=1} \geq 0.85$ and $\hat{\sigma}$ -bin gap ≤ 0.20 on both datasets** (Tab. 9, Tab. 10; all rows $B = 200$ chrom-bootstrap mean $\pm$ std). The on-axis B2 $\hat{\sigma}$ -Mondrian [Boström and Johansson, 2020] attains local coverage shape but loses minority coverage ($\text{cov}_{|Y=1} = 0.624$ on Complex), so its $\hat{\sigma}$ -bin gap is dominated by minority-cell undercoverage; in this setting HCCP’s joint partition keeps both within the target band (HCCP gap 0.060/0.173 Complex/Mendelian vs. B2 0.827/0.401, bootstrap means; B3 class-Mondrian [Sadinle et al., 2019] parallel; per-baseline ratios in Tab. 9 caption). The Mendelian advantage narrows under weighted CP at the recommended operating point (HCCP 0.173 ± 0.192 vs. wCP 0.252 ± 0.153 , paired-bootstrap 74.5% wins, CI crosses zero); a $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$ sweep (App. C.6) shows HCCP holds the target band against all baselines across the full range on Complex but reverses on Mendelian at $K_{\text{eval}} \geq 5$ (per-cell minority < 70). The reversal traces directly to Eq. (3): below the per-cell variance-stable floor, the $K/(\pi_{\min}n)$ variance term overtakes the $L_F R/K$ Lipschitz bias gain, and the smoother reweighting of weighted CP avoids the hard-bin variance penalty. This is a regime boundary, not a model-class limitation. The same calibration on the published LogReg aggregator yields essentially identical $\hat{\sigma}$ -bin gap (App. C.5), decoupling the conformal-side gain from base-predictor accuracy. The Pareto-uniqueness claim is not knife-edge: HCCP is the unique frontier on the full $\{0.80, 0.85, 0.90\} \times \{0.15, 0.20, 0.25, 0.30\}$ grid on Complex and on a $\{0.80, 0.85\} \times \{0.20, 0.25\}$ rectangle on Mendelian; the operating point $(0.85, 0.20)$ is interior, not on the boundary (App. C.4).

Imbalance scaling validates T5’s $\pi_{\min}$ dependence. Subsampling Complex to $\pi_+ \in \{0.05, 0.10, 0.20, 0.30, 0.50\}$, B1 and B2 minority-coverage degrades monotonically as $\pi_+ \rightarrow 0$ ($0.587 \rightarrow 0.810$ for B1; $0.650 \rightarrow 0.878$ for B2), while HCCP and class-Mondrian hold at the target across the full sweep (App. C, Tab. 11) — the empirical realisation of the $\pi_{\min}^{-1/2}$ factor in T5’s bound.

6.3 Three-axis OOD + cross-domain stress test

We evaluate HCCP on three orthogonal TraitGym distributional axes (chrom-LOO, trait-LOO, cross-dataset) and two external benchmarks (ProteinGym cross-domain, Open Targets cross-platform).

TraitGym 3-axis OOD. Trait-LOO reaches the finite-sample $1/(n_{k,b} + 1)$ floor on both datasets ($\hat{\sigma}$ -bin gap ≤ 0.004), with a single exception: the Skeletal/connective OMIM cluster ($n_{\text{tot}} = 470$) shows $\text{cov}_{|Y=1} = 0.745$ (1/3 Mendelian clusters fall below the 0.85 minority-coverage target). *Chrom-LOO* on the headline CADD+GPN-MSA+Borzoï feature set: per-chrom lcoverage gap ≤ 0.100 Mendelian (chr6 at 0.800) and ≤ 0.023 Complex (Tab. 16). *Cross-dataset* C $\rightarrow$ M holds (gap 0.035); the reverse M $\rightarrow$ C direction is a joint-failure case study (§7, R3).

External benchmarks. ProteinGym cross-domain [Notin et al., 2023]: across 50 assay-LOO held-out proteins, HCCP attains marginal coverage 0.887, within ± 0.02 of target on 27/50 assays and within ± 0.05 on 44/50. The 6 outlier assays correlate with low base-predictor AUPRC (Spearman $\rho = 0.46$, $p = 0.001$), confirming coverage failures track base-predictor weakness rather than conformal failure (App. E). *Open Targets cross-platform* ($n = 11,400$, $\pi_+ = 0.10$, feature pipeline disjoint from TraitGym; App. G): HCCP attains $\hat{\sigma}$ -bin gap 0.054 vs. 0.900 for all three matched- $(\hat{p}, \hat{\sigma})$ baselines ($16.7\times$, 93% of paired chrom-resamples favour HCCP, $p_{\text{one}} = 0.075$); the K_{eval} reversal at per-cell-minority < 120 replicates the TraitGym Mendelian boundary on a feature-pipeline-disjoint benchmark (Tab. 28, Tab. 29).

6.4 Ablations

Ablations on K -sweep (T5 rate validation), $\hat{\sigma}$ source (supervised vs. DEGU-lite ensemble), $\hat{\sigma}$ target function (residual vs. log-variance), nonconformity score, class-conditional on/off, AlphaGenome [Avsec et al., 2026] feature extension, and synthetic n -sweep (verifying the $O(n^{-1/2})$ rate of Theorem 3) are reported in App. C and App. F. The synthetic sweep is the direct test of T5.1’s dimension-free claim: log-log slope -0.44 ± 0.08 across $d \in \{8, 32, 128\}$ and $\pi_{\min} \in \{0.1, 0.5\}$ (T5.1 prediction -0.50 within 1 SD); the empirical $\bar{C}_{0.10}/\bar{C}_{0.50} = 1.62$ is below the $\sqrt{5} \approx 2.24$

asymptotic ratio (T5.1 is an upper bound on G , so empirical ratios $\leq \sqrt{5}$ are consistent with the bound). The DEGU-lite ablation [Zhou et al., 2026] swaps our supervised residual $\hat{\sigma}$ for an ensemble-distilled Gaussian-uncertainty $\hat{\sigma}$ at matched GBM backbone and matched HCCP-style Mondrian calibration; this isolates the choice of variance source within HCCP and is *not* a head-to-head against DEGU’s published Keras CNN MPRA pipeline. With this caveat, the supervised residual head outperforms the distilled ensemble head by $4.3\times$ on Complex worst-cell coverage, identifying where aleatoric-aware $\hat{\sigma}$ binds (polygenic complex traits) and where ensemble disagreement suffices (monogenic Mendelian).

7 Discussion

Clinical implications. HCCP prediction sets are not a replacement for expert review; they are a calibrated abstention signal. A singleton set $\mathcal{C}_\alpha(x) = \{k\}$ is a confident call; a set containing both labels or neither is an explicit abstention. On TraitGym Mendelian, HCCP produces singletons for $\approx 78\%$ of variants at $\alpha = 0.10$ ($K_{\text{eval}} = 5$, Tab. 2; cf. 0.827 at $K_{\text{eval}} = 3$ in Tab. 10). Operating-point selection by miscoverage level is dataset-specific: $\alpha = 0.10$ for high-stakes rare-disease classification, $\alpha = 0.20$ for population-screening triage of complex traits where the abstention burden is otherwise prohibitive (cf. Adegoke et al. [2025] report 41% abstention at $\alpha \approx 0.20$ in a clinical screening setting); App. D reports the singleton-vs-abstention tradeoff across $\alpha \in [0.05, 0.50]$.

Operational envelope and limitations. The coverage guarantees of HCCP hold when (a) per-cell minority ≥ 100 , (b) base AUPRC $\gtrsim 0.30$, and (c) per-cell KS-reject $\lesssim 50\%$. Assumption 1 (chrom-group exchangeability) is load-bearing for T3; T3’ absorbs A2-cell violations via the chrom-TV term on Mendelian (46.2% reject, K -invariant in $[40, 46]\%$ across $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$; App. B); Assumption 4 is required only for T5. We report four labelled out-of-envelope regimes as case studies, not coverage claims: **(R1)** Mendelian $K_{\text{eval}} \geq 5$ violates (a) at per-cell minority < 70 (§6.2; reversal *anticipated* by Eq. (3)’s variance-term scaling); **(R2)** the Skeletal/connective OMIM cluster ($n = 470$, $1/3$ Mendelian clusters) violates (a) with $\text{cov}_{|Y=1} = 0.745$ (§6.3); **(R3)** the M→C cross-dataset direction jointly violates (b) (AUPRC 0.205 vs. ID 0.353) and (c) (ClinVar vs. GWAS catalog selection [Tibshirani et al., 2019]), with marginal coverage 0.738 (0.004 below the Barber 2023 Thm. 2 proxy bound of 0.742); the forward C→M and ProteinGym cross-domain violate neither precondition and HCCP holds; **(R4)** ProteinGym 6/50 outliers violate (b) (CD19 AUPRC 0.24). Additional scope conditions (TV proxy approximation, large-sample K^* vs. finite-sample per-cell floor binding, DEGU CNN end-to-end port out of scope, per-assay $\hat{\sigma}$ adaptation as future work) are detailed in App. B, C, E.

Broader impact. VEP systems inform consequential clinical decisions (cascade testing, reproductive choices). HCCP provides calibrated abstention but does not guarantee clinical safety: the $\approx 78\%$ singletons on Mendelian still carry a 10% finite-sample miscoverage floor, and abstention must be paired with expert review to translate coverage into clinical utility. Population-stratified deployment requires re-calibration on ancestry-conditional folds (a fourth Mondrian axis); TraitGym is built from ClinVar and the GWAS catalog and over-represents European-ancestry individuals. The ACMG/AMP consensus [Richards et al., 2015] treats computational predictions as PP3/BP4 supporting evidence, not stand-alone classification; HCCP singletons should be integrated into multi-evidence frameworks, not deployed as diagnostic.

8 Conclusion

We introduced HCCP, attaining class-conditional and bin-local conformal coverage jointly under imbalance and heteroscedasticity, with an $O(n^{-1/2})$ finite-sample rate within the equi-bin Mondrian- K class matched by a within-class $\Omega(n^{-1/2})$ lower bound and $\pi_{\min}^{-1/2}$ explicit in the constant (T5.1, T5.2; Theorems 3, 4). On TraitGym, Open Targets matched-9, ProteinGym, and synthetic data, HCCP holds the target band on Complex and competes with weighted CP on Mendelian within a $K_{\text{eval}} = 3$ window — itself a regime boundary predicted by Eq. 3, not a model-class limitation.

References

- Adele Alex Adegoke, Idris Babalola, and Peter Adebayo Odesola. A step toward the use of AI for polycystic ovary syndrome (PCOS): Staged modelling with uncertainty-aware triage and conformal prediction with cost-efficient risk. *medRxiv*, 2025. doi: 10.1101/2025.10.27.25338898.
- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *Foundations and Trends in Machine Learning*, 16(4), 2022.
- Žiga Avsec, Natasha Latysheva, Jun Cheng, Guido Novati, Kyle R. Taylor, Tom Ward, Clare Bycroft, Lauren Nicolaisen, Eirini Arvaniti, Joshua Pan, Demis Hassabis, and Pushmeet Kohli. Advancing regulatory variant effect prediction with AlphaGenome. *Nature*, 649:1206–1218, 2026. doi: 10.1038/s41586-025-10014-0.
- Konstantina Bairaktari, Jiayun Wu, and Zhiwei Steven Wu. Kandinsky conformal prediction: Beyond class- and covariate-conditional coverage. In *ICML*, 2025.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. The limits of distribution-free conditional predictive inference. *Information and Inference*, 2020.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Predictive inference with the jackknife+. *Annals of Statistics*, 49(1):486–507, 2021.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Conformal prediction beyond exchangeability. *Annals of Statistics*, 51(2):816–845, 2023.
- Gonzalo Benegas, Carlos Albors, Alan J. Aw, Chengzhong Ye, and Yun S. Song. GPN-MSA: Unsupervised variant effect prediction with genome-wide MSA. *bioRxiv*, 2023.
- Gonzalo Benegas, Gökçen Eraslan, and Yun S. Song. TraitGym: A benchmark for deep learning on complex genomic variants. *bioRxiv*, 2025. doi: 10.1101/2025.04.07.646951.
- Henrik Boström. crepes: a python package for generating conformal regressors and predictive systems. In *Proceedings of the Eleventh Symposium on Conformal and Probabilistic Prediction with Applications (COPA)*, PMLR 179, pages 24–41, 2022.
- Henrik Boström and Ulf Johansson. Mondrian conformal regressors. In *Proceedings of the Ninth Symposium on Conformal and Probabilistic Prediction with Applications (COPA)*, PMLR 128, pages 114–133, 2020.
- Robin De Wolf, Bernard De Baets, and Willem Waegeman. Conditional validity of heteroskedastic conformal regression. *Information and Inference: A Journal of the IMA*, 2025.
- Hang Deng, Qiyang Han, and Cun-Hui Zhang. Confidence intervals for the means of the selected populations and convergence rates of isotonic estimators. *arXiv preprint arXiv:2103.04340*, 2021.
- Tiffany Ding, Anastasios N. Angelopoulos, Stephen Bates, Michael I. Jordan, and Ryan J. Tibshirani. Class-conditional conformal prediction with many classes. In *NeurIPS*, 2023.
- Isaac Gibbs and Emmanuel J. Candès. Adaptive conformal inference under distribution shift. In *NeurIPS*, 2021.
- Rohan Hore and Rina Foygel Barber. Conformal prediction with local weights: randomization enables robust guarantees. *Journal of the Royal Statistical Society: Series B*, 2025.
- Julien Walden Huang, Stephen J. Roberts, and Jan-Peter Calliess. On the sample complexity of lipschitz constant estimation. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=UIa1YAHdBH>. ICLR 2024 Journal Track.
- Martin Kircher, Daniela M. Witten, Preti Jain, Brian J. O’Roak, Gregory M. Cooper, and Jay Shendure. A general framework for estimating the relative pathogenicity of human genetic variants. *Nature Genetics*, 2014.

- 408 Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive
409 uncertainty estimation using deep ensembles. In *NeurIPS*, 2017.
- 410 Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-
411 free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):
412 1094–1111, 2018.
- 413 Johannes Linder, Divyanshi Srivastava, Han Yuan, Vikram Agarwal, and David R. Kelley. Predicting
414 RNA-seq coverage from DNA sequence as a unifying model of gene regulation. *Nature Genetics*,
415 2025.
- 416 David A. Nix and Andreas S. Weigend. Estimating the mean and variance of the target probability
417 distribution. In *IEEE ICNN*, 1994.
- 418 Pascal Notin, Aaron Kollasch, Daniel Ritter, Lood van Niekerk, Steffanie Paul, Han Spinner, Nathan
419 Rollins, Ada Shaw, Rose Orenbuch, Ruben Weitzman, Jonathan Frazer, Mafalda Dias, Dinko
420 Franceschi, Yarin Gal, and Debora Marks. ProteinGym: Large-scale benchmarks for protein fitness
421 prediction and design. In *NeurIPS Datasets and Benchmarks*, 2023.
- 422 David Ochoa, Andrew Hercules, Miguel Carmona, Daniel Suveges, et al. The next-generation open
423 targets platform: reimaged, redesigned, rebuilt. *Nucleic Acids Research*, 51(D1):D1353–D1359,
424 2023.
- 425 Harris Papadopoulos, Alex Gammerman, and Vladimir Vovk. Normalized nonconformity measures
426 for regression conformal prediction. In *Proceedings of the IASTED International Conference on*
427 *Artificial Intelligence and Applications (AIA)*, pages 64–69. ACTA Press, 2008.
- 428 Christina Papangelou, Konstantinos Kyriakidis, Pantelis Natsiavas, Ioanna Chouvarda, and Andigoni
429 Malousi. Reliable machine learning models in genomic medicine using conformal prediction.
430 *Frontiers in Bioinformatics*, 2025. doi: 10.3389/fbinf.2025.1507448.
- 431 Vincent Plessier, Alexander Fishkov, Mohsen Guizani, Maxim Panov, and Eric Moulines. Probabilistic
432 conformal prediction with approximate conditional validity. In *ICLR*, 2025.
- 433 Sue Richards, Nazneen Aziz, Sherri Bale, David Bick, Soma Das, Julie Gastier-Foster, Wayne W.
434 Grody, Madhuri Hegde, Elaine Lyon, Elaine Spector, Karl Voelkerding, and Heidi L. Rehm. Stan-
435 dards and guidelines for the interpretation of sequence variants: a joint consensus recommendation
436 of the American College of Medical Genetics and Genomics and the Association for Molecular
437 Pathology. *Genetics in Medicine*, 17(5):405–424, 2015.
- 438 Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. In
439 *NeurIPS*, 2019.
- 440 Mauricio Sadinle, Jing Lei, and Larry Wasserman. Least ambiguous set-valued classifiers with
441 bounded error levels. *Journal of the American Statistical Association*, 114(525):223–234, 2019.
- 442 Yuanjie Shi, Subhankar Ghosh, Taha Belkhouja, Janardhan Rao Doppa, and Yan Yan. Conformal
443 prediction for class-wise coverage via augmented label rank calibration. In *NeurIPS*, 2024.
- 444 Ryan J. Tibshirani, Rina Foygel Barber, Emmanuel J. Candès, and Aaditya Ramdas. Conformal
445 prediction under covariate shift. In *NeurIPS*, 2019.
- 446 Paolo Toccaceli. CRPS-optimal binning for univariate conformal regression. *arXiv preprint*
447 *arXiv:2603.22000*, 2026.
- 448 Lars van der Laan. Self-calibrating conformal prediction. In *NeurIPS*, 2024.
- 449 Vladimir Vovk, David Lindsay, Ilia Nouretdinov, and Alex Gammerman. Mondrian confidence
450 machine. Technical report, Royal Holloway, University of London, 2003.
- 451 Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random*
452 *World*. Springer, 2005.

- Chen Wang, Fan Wang, Malgorzata Bogdan, Marco Masala, Edoardo Fiorillo, Marcella Devoto, Francesco Cucca, Dan Belsky, and Iuliana Ionita-Laza. Individualized uncertainty quantification in polygenic risk scores using conformalized quantile regression. *bioRxiv*, 2025. doi: 10.1101/2025.10.13.682102.
- Tianmin Xie, Yanfei Zhou, Ziyi Liang, Stefano Favaro, and Matteo Sesia. Conformal inference for open-set and imbalanced classification. *arXiv preprint arXiv:2510.13037*, 2025.
- Yunzhen Yao, Lie He, and Michael Gastpar. Non-asymptotic analysis of efficiency in conformalized regression. In *ICLR*, 2026. arXiv:2510.07093.
- Jessica Zhou, Kaeli Rizzo, Trevor Christensen, Ziqi Tang, and Peter K. Koo. Uncertainty-aware genomic deep learning with knowledge distillation. *npj Artificial Intelligence*, 2(3), 2026. URL <https://www.nature.com/articles/s44387-025-00053-3>.

A Full proofs of T1–T5

A.1 Assumptions hierarchy

We collect in one place the assumptions invoked across §5. The chain of A2 variants is:

$$\underbrace{\text{A2-cell}}_{\text{T3}} \implies \underbrace{\text{A2-class}}_{\text{T2}} \implies \underbrace{\text{A2-marginal}}_{\text{T1}},$$

where each level is obtained from the stronger one by marginalizing over a partition variable.

- **A1** (chrom-group exchangeability, 1): conditional on $C = c$, $\{(X_i, Y_i) : C_i = c\}$ are exchangeable with law P_c ; different chromosomes are mutually independent.
- **A1'** (chrom-wise i.i.d., 2): the stronger version in which $\{(X_i, Y_i) : C_i = c\}$ are i.i.d. P_c .
- **A2-marginal**: $S \mid C = c \stackrel{d}{=} S \mid C = c'$ for all $c, c' \in \mathcal{C}$; the common law is F .
- **A2-class**: $S \mid (Y = k, C = c) \stackrel{d}{=} S \mid (Y = k, C = c')$ for all $c, c', k \in \{0, 1\}$; common law F_k .
- **A2-cell**: $S \mid (Y = k, b(X) = b, C = c) \stackrel{d}{=} S \mid (Y = k, b(X) = b, C = c')$ for all c, c', k, b ; common law F_{kb} .
- **A3-loc** (local stationarity, used in T3-loc): for every (x_0, r) satisfying the T3-loc resolution condition, $s(X, k) \mid (Y = k, b(X) = b, X \in B_\ell(x_0, r))$ is distributed as F_{kb} .
- **B1** ($\hat{\sigma}$ relative error, used in T3.b): $\|\hat{\sigma} - \sigma^*\|_\infty / \inf_x \sigma^*(x) \leq \eta$ for some $\eta \in (0, 1)$.
- **A-SL** (score-Lipschitz regularity, 4, used in T5): $\sup_s |F_{k, \sigma_1}(s) - F_{k, \sigma_2}(s)| \leq L_F \cdot |\sigma_1 - \sigma_2|$ for all $\sigma_1, \sigma_2 \in [\hat{\sigma}_{\min}, \hat{\sigma}_{\max}]$ and $k \in \{0, 1\}$.

Empirical validation of A2. The within-cell KS audit (Tab. 4) ran 177 cell-level two-sample KS tests across both datasets ($K_{\text{eval}} = 5 \text{ bins} \times 2 \text{ classes} \times \text{test chromosomes with non-empty cells}$; 52 tests on Mendelian, 125 on Complex). The rejection rates at nominal $p < 0.05$ are 46.2% (Mendelian, max KS 0.434) and 10.4% (Complex, max KS 0.315). Since $d_{\text{TV}} \leq d_{\text{KS}}/\sqrt{2}$ does *not* hold in general — only $d_{\text{TV}} \leq d_{\text{KS}}$ for univariate F , with the tighter $\sqrt{2}$ bound under additional regularity — we use the conservative envelope $\bar{\delta}_{\text{TV}}^{(k,b)} \leq d_{\text{KS}}$, giving worst-case $\bar{\delta}_{\text{TV}}^{(k,b)} \leq 0.31$ on the binding Mendelian cell (chr9, chr7, chr3 pathogenic-rich at the highest- $\hat{\sigma}$ bin). A2-cell holds approximately on Complex; T3' supplies the certificate on Mendelian.

A.2 Auxiliary lemmas

We state four standard lemmas used in the proofs below.

Lemma 1 (Rank uniformity for exchangeable sequences). *Let $T_1, \dots, T_{m+1}$ be exchangeable random variables with a continuous joint distribution. Then the rank of T_{m+1} among $(T_1, \dots, T_{m+1})$ is uniformly distributed on $\{1, \dots, m+1\}$.*

494 *Proof.* By exchangeability, for any permutation π of $\{1, \dots, m+1\}$, the joint law of
 495 $(T_{\pi(1)}, \dots, T_{\pi(m+1)})$ equals that of $(T_1, \dots, T_{m+1})$. Continuity implies $P(T_i = T_j) = 0$ for $i \neq j$,
 496 so the ranking is almost surely unique. For each $r \in \{1, \dots, m+1\}$, the event $\{\text{rank}(T_{m+1}) = r\}$
 497 can be mapped to the event $\{\text{rank}(T_{\pi(m+1)}) = r\}$ by a suitable permutation that places the element
 498 currently at rank r in position $m+1$, so all rank values are equiprobable. $\square$

499 **Lemma 2** (Exchangeable CP quantile bound). *Under the conditions of Lemma 1, let $\hat{q} =$
 500 $\text{Quantile}(\{T_1, \dots, T_m\}, \lceil (m+1)(1-\alpha) \rceil / m)$ be the $\lceil (m+1)(1-\alpha) \rceil$ -th order statistic of the
 501 first m values. Then*

$$P(T_{m+1} \leq \hat{q}) = \frac{\lceil (m+1)(1-\alpha) \rceil}{m+1} \in \left[1-\alpha, 1-\alpha + \frac{1}{m+1} \right].$$

502 *Proof.* By Lemma 1, T_{m+1} is at rank r with probability $1/(m+1)$ for each r . The event $\{T_{m+1} \leq \hat{q}\}$
 503 is exactly $\{\text{rank}(T_{m+1}) \leq \lceil (m+1)(1-\alpha) \rceil\}$ (since $\hat{q}$ is the $\lceil (m+1)(1-\alpha) \rceil$ -th order statistic of
 504 $(T_1, \dots, T_m)$, which is also the $\lceil (m+1)(1-\alpha) \rceil$ -th order statistic of the augmented set when T_{m+1}
 505 has rank at most that value). Thus

$$P(T_{m+1} \leq \hat{q}) = \frac{\lceil (m+1)(1-\alpha) \rceil}{m+1}.$$

506 Since $\lceil (m+1)(1-\alpha) \rceil \in [(m+1)(1-\alpha), (m+1)(1-\alpha) + 1]$, dividing by $m+1$ gives the
 507 stated bounds. $\square$

508 **Lemma 3** (Monotone score-rescaling invariance). *Let $g : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ be a strictly increasing function.
 509 If $\hat{q}$ is the conformal quantile computed from scores $\{S_i\}$ and $\hat{q}' = g(\hat{q})$ is the quantile computed
 510 from rescaled scores $\{g(S_i)\}$, then for any test score S_{test} ,*

$$S_{\text{test}} \leq \hat{q} \iff g(S_{\text{test}}) \leq \hat{q}'.$$

511 *Proof.* The quantile index $\lceil (m+1)(1-\alpha) \rceil$ is the same for both sets. Since g is strictly increasing,
 512 it preserves the ordering of scores. Thus the j -th order statistic of $\{g(S_i)\}$ equals g applied to the j -th
 513 order statistic of $\{S_i\}$, giving $\hat{q}' = g(\hat{q})$. The equivalence follows from g being order-preserving. $\square$

514 **Lemma 4** (Bin-boundary mass bound). *For equiprobable bins with K bins and minimum bin width
 515 Δ , the probability mass of $\sigma^*(X)$ within distance δ of any bin edge is at most $2K\delta/\Delta$.*

516 *Proof.* Each bin contains mass $1/K$ and has width $\geq \Delta$, so the density of $\sigma^*(X)$ is at most
 517 $1/(K\Delta) \cdot K = 1/\Delta$ on average. There are $K-1$ internal bin edges, and each has a neighborhood
 518 of width 2δ that captures mass at most $2\delta/\Delta$. Summing over all edges gives $2(K-1)\delta/\Delta \leq$
 519 $2K\delta/\Delta$. $\square$

520 A.3 Lemma 5: T1 marginal coverage

521 **Lemma 5** (T1 marginal coverage — standard split-CP at the marginal partition level). *Under 2 and
 522 A2, for test point $(X, Y) \sim P_{c^*}$ independent of the calibration fold and continuous score law,*

$$1-\alpha \leq P(Y \in \mathcal{C}_\alpha(X)) \leq 1-\alpha + \frac{1}{n_{\min} + 1}, \quad n_{\min} = \min_k n_k.$$

523 *Proof.* We decompose the marginal coverage via the true label.

524 *Step 1 (Reduction to per-class coverage).* By the law of total probability:

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})) = \sum_{k \in \{0,1\}} P(Y_{\text{test}} = k) \cdot P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) \mid Y_{\text{test}} = k).$$

525 Since $Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})$ reduces to $s(X_{\text{test}}, Y_{\text{test}}) \leq \hat{q}_{Y_{\text{test}}}$ by definition of $\mathcal{C}_\alpha$, we have

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) \mid Y_{\text{test}} = k) = P(s(X_{\text{test}}, k) \leq \hat{q}_k \mid Y_{\text{test}} = k).$$

526 *Step 2 (Joint exchangeability inside class k).* Fix k . Under 2, for each calibration chromosome $c \neq c^*$,
 527 $(X_i, Y_i)_{C_i=c}$ are i.i.d. P_c . The score $S_i = s(X_i, Y_i)$ is a deterministic function of (X_i, Y_i) via the

fixed $(\hat{p}, \hat{\sigma})$ trained on $\mathcal{D}^{(-c^*)}$. Restricting to $i \in \mathcal{I}_k^{(-c^*)}$ (i.e., $Y_i = k$) yields i.i.d. samples from a chrom-specific conditional law F_k^c . By A2-class (implied by A2-marginal when class prevalence $P(Y = k | C)$ is chrom-independent, or invoked directly), $F_k^c = F_k$ for all c . Across chromosomes, samples are mutually independent (A1). Hence the full calibration set $\{S_i\}_{i \in \mathcal{I}_k^{(-c^*)}}$ is i.i.d. F_k .

Conditional on $Y_{\text{test}} = k$, the test score $s(X_{\text{test}}, k)$ has distribution F_k (by A2-class, with $c = c^*$) and is independent of the calibration samples (different chromosome, A1). Thus the augmented sequence $(S_1^{(k)}, \dots, S_{n_k}^{(k)}, s(X_{\text{test}}, k))$ is i.i.d. F_k , hence exchangeable.

Step 3 (Standard exchangeable CP bound). By Lemma 2 applied with $m = n_k$ and $T_{m+1} = s(X_{\text{test}}, k)$:

$$P(s(X_{\text{test}}, k) \leq \hat{q}_k | Y_{\text{test}} = k) = \frac{\lceil (n_k + 1)(1 - \alpha) \rceil}{n_k + 1} \in \left[1 - \alpha, 1 - \alpha + \frac{1}{n_k + 1} \right].$$

Step 4 (Assemble). Substituting into Step 1:

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})) = \sum_k P(Y_{\text{test}} = k) \cdot P(s(X_{\text{test}}, k) \leq \hat{q}_k | Y_{\text{test}} = k).$$

Each term lies in $[1 - \alpha, 1 - \alpha + 1/(n_k + 1)]$. The convex combination inherits the range: the lower bound gives $\geq 1 - \alpha$; the upper bound gives $\leq 1 - \alpha + 1/(n_{\min} + 1)$ with $n_{\min} = \min_k n_k$. $\square$

A.4 Lemma 6: T2 class-conditional coverage

Lemma 6 (T2 class-conditional coverage — standard split-CP at the class partition level). *Under the conditions of Lemma 5 with A2 at class level and $n_k \geq \lceil 1/\alpha \rceil$ for each $k \in \{0, 1\}$,*

$$1 - \alpha \leq P(Y \in \mathcal{C}_\alpha(X) | Y = k) \leq 1 - \alpha + \frac{1}{n_k + 1}.$$

Proof. Fix $k \in \{0, 1\}$ with $n_k \geq \lceil 1/\alpha \rceil$.

Step 1 (Exchangeability). Under 2 and A2-class, conditioning on $Y_i = k$ makes $(S_1^{(k)}, \dots, S_{n_k}^{(k)})$ i.i.d. F_k by the same argument as T1, Step 2. Conditional on $Y_{\text{test}} = k$, $s(X_{\text{test}}, k)$ is independent of calibration and distributed as F_k . Thus the joint $(S_1^{(k)}, \dots, S_{n_k}^{(k)}, s(X_{\text{test}}, k))$ is i.i.d., hence exchangeable.

Step 2 (Quantile bound). The prediction-set inclusion event $\{Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})\}$ under $Y_{\text{test}} = k$ is $\{s(X_{\text{test}}, k) \leq \hat{q}_k\}$. By Lemma 2 with $m = n_k$:

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) | Y_{\text{test}} = k) = \frac{\lceil (n_k + 1)(1 - \alpha) \rceil}{n_k + 1} \in \left[1 - \alpha, 1 - \alpha + \frac{1}{n_k + 1} \right].$$

The sample-size requirement $n_k \geq \lceil 1/\alpha \rceil$ ensures $\lceil (n_k + 1)(1 - \alpha) \rceil \leq n_k$, so $\hat{q}_k < +\infty$ and the prediction set is non-trivial. $\square$

Remark. T1 is the law-of-total-probability average of the two T2 per-class bounds, so T2 is strictly stronger. The sample-size requirement is comfortably met on TraitGym: at $\alpha = 0.10$ we need $n_k \geq 10$; the smallest chrom-LOO calibration has $n_1 \geq 30$ (Mendelian) and $n_1 \geq 100$ (Complex).

A.5 Robust fallback: T1' and T2'

When A2 fails (e.g., a test chromosome has systematically different variant characteristics), we invoke the following restatement of Barber et al. [2023] Theorem 2 in our notation.

Barber 2023 Theorem 2 (restated). Let $Z = (Z_1, \dots, Z_n, Z_{n+1})$ be the full data ordering with Z_{n+1} as the test point, $R(Z) = (S_1, \dots, S_{n+1})$ the residual vector, and $R(Z^{(i)})$ the sequence with Z_i and Z_{n+1} swapped. For prespecified weights $w_i \in [0, 1]$ with $\tilde{w}_i = w_i/(w_1 + \dots + w_n + 1)$,

$$P(Y_{n+1} \in \hat{C}_n(X_{n+1})) \geq 1 - \alpha - \sum_{i=1}^n \tilde{w}_i \cdot d_{\text{TV}}(R(Z), R(Z^{(i)})).$$

No factor of 2.

562 *Proof of T1' and T2'.* Use uniform weights $w_i = 1$, so $\tilde{w}_i = 1/(n+1)$. The Barber 2023 bound
 563 becomes

$$P(\text{coverage}) \geq 1 - \alpha - \frac{1}{n+1} \sum_i d_{\text{TV}}(R(Z), R(Z^{(i)})).$$

564 Under 1 alone (without A2), $R(Z^{(i)})$ differs from $R(Z)$ in the position of one residual: $S_i \leftrightarrow S_{\text{test}}$.
 565 If S_i is from chromosome c (so $S_i \sim F_c$) and S_{test} from c^* (so $S_{\text{test}} \sim F_{c^*}$), the swap changes
 566 the marginal law of exactly two entries of the residual vector. Under A1 (cross-chromosome
 567 independence), entries from different chromosomes are independent, so the TV between $R(Z)$ and
 568 $R(Z^{(i)})$ is bounded by the TV between the two swapped marginals:

$$d_{\text{TV}}(R(Z), R(Z^{(i)})) \leq d_{\text{TV}}(F_c, F_{c^*}).$$

569 Substituting and grouping calibration points by chromosome (each of the $|\mathcal{C}| - 1$ calibration chromo-
 570 somes contributes equally to the average under uniform weights):

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})) \geq 1 - \alpha - \bar{\delta}_{\text{TV}}, \quad \bar{\delta}_{\text{TV}} = \frac{1}{|\mathcal{C}| - 1} \sum_{c \neq c^*} d_{\text{TV}}(F_c, F_{c^*}).$$

571 This is T1'.

572 For T2', the same argument is applied per class: restrict to calibration indices with $Y_i = k$, replace
 573 F_c with $F_c^{(k)}$, and obtain

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) \mid Y_{\text{test}} = k) \geq 1 - \alpha - \bar{\delta}_{\text{TV}}^{(k)}, \quad \bar{\delta}_{\text{TV}}^{(k)} = \frac{1}{|\mathcal{C}| - 1} \sum_{c \neq c^*} d_{\text{TV}}(F_c^{(k)}, F_{c^*}^{(k)}).$$

574 Both bounds carry no factor of 2 (weighted-swap Barber 2023, not jackknife+). $\square$

575 **Empirical TV estimates.** The cell-level KS audit (App. B.4, 177 tests across both datasets) gives
 576 worst-cell $d_{\text{KS}} = 0.434$ on Mendelian (chr9 / chr7 / chr3 highest- $\hat{\sigma}$ pathogenic-rich cell) and 0.315
 577 on Complex. Using $\bar{\delta}_{\text{TV}}^{(k,b)} \leq d_{\text{KS}}$ as the conservative envelope, the cell-level coverage gap is bounded
 578 by ≤ 0.31 on Mendelian and ≤ 0.05 on Complex; observed per-chromosome coverage gaps remain
 579 ≤ 0.100 throughout (chr6 at 0.800; full table in App. C.7), well below the certificate.

580 A.6 Proof of Theorem 1 (bin-conditional coverage, exact)

581 *Proof.* The argument is the class-conditional proof of Lemma 6 applied inside a single Mondrian cell
 582 (k, b) , with A2-cell replacing A2-class. Fix a cell (k, b) with $n_{kb} \geq \lceil 1/\alpha \rceil$.

583 *Step 1 (Cell restriction).* The event “ $Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}})$ given $Y_{\text{test}} = k$ ” is exactly $\{s(X_{\text{test}}, k) \leq$
 584 $\hat{q}_{k,b}(X_{\text{test}})\}$ by the definition of $\mathcal{C}_\alpha$. Conditioning further on $b(X_{\text{test}}) = b$ fixes the threshold to $\hat{q}_{kb}$.

585 *Step 2 (I.i.d. F_{kb} via A1' + A2-cell).* Under 2, for each calibration chromosome $c \neq c^*$, $(X_i, Y_i)_{C_i=c}$
 586 are i.i.d. P_c ; the score $S_i = s(X_i, Y_i)$ is a deterministic function of (X_i, Y_i) via the fixed $(\hat{p}, \hat{\sigma})$.
 587 Restricting to $i \in \mathcal{I}_{kb}^{(-c^*)}$ (i.e., $Y_i = k$ and $b(X_i) = b$) yields i.i.d. samples from a chromosome-
 588 specific conditional law, which by A2-cell equals F_{kb} for every calibration chromosome. Across
 589 chromosomes, samples are mutually independent (1). Hence the full calibration set $\{S_i\}_{i \in \mathcal{I}_{kb}^{(-c^*)}}$ is
 590 i.i.d. F_{kb} .

591 *Step 3 (Test-score independence and exchangeability).* Conditional on $Y_{\text{test}} = k$ and $b(X_{\text{test}}) = b$,
 592 the test score $s(X_{\text{test}}, k)$ has distribution F_{kb} (by A2-cell, with $c = c^*$) and is independent of the
 593 calibration samples (different chromosome, 1). Thus $(S_1^{(kb)}, \dots, S_{n_{kb}}^{(kb)}, s(X_{\text{test}}, k))$ is i.i.d. F_{kb} , in
 594 particular exchangeable.

595 *Step 4 (CP quantile bound).* By Lemma 2 with $m = n_{kb}$:

$$P(s(X_{\text{test}}, k) \leq \hat{q}_{kb}) = \frac{\lceil (n_{kb} + 1)(1 - \alpha) \rceil}{n_{kb} + 1} \in \left[1 - \alpha, 1 - \alpha + \frac{1}{n_{kb} + 1} \right].$$

596 Combining Steps 1 and 4 yields

$$1 - \alpha \leq P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) \mid Y_{\text{test}} = k, b(X_{\text{test}}) = b) \leq 1 - \alpha + \frac{1}{n_{kb} + 1}. \quad \square$$

597 **Local coverage gap bound.** For each cell (k, b) ,

$$|P(Y \in \mathcal{C}_\alpha(X) \mid Y = k, b(X) = b) - (1 - \alpha)| \leq \frac{1}{n_{kb} + 1}.$$

598 **A.7 Lemma 7: T3' bin-conditional coverage, robust**

599 **Lemma 7** (T3' bin-conditional coverage, robust). *Under 1 alone (no A2 at any level), for each cell*
600 *(k, b) with $n_{k,b} \geq \lceil 1/\alpha \rceil$,*

$$P(Y \in \mathcal{C}_\alpha(X) \mid Y = k, b(X) = b) \geq 1 - \alpha - \bar{\delta}_{\text{TV}}^{(k,b)},$$

$$\bar{\delta}_{\text{TV}}^{(k,b)} = \frac{1}{|\mathcal{C}| - 1} \sum_{c \neq c^*} d_{\text{TV}}(F_c^{(k,b)}, F_{c^*}^{(k,b)}).$$

601 *Proof.* Apply Barber et al. [2023] Theorem 2 with calibration weights

$$w_i = \mathbb{1}\{i \in \mathcal{I}_{kb}^{(-c^*)}\}, \quad \tilde{w}_i = \frac{w_i}{w_1 + \dots + w_n + 1}.$$

602 So $\tilde{w}_i = 1/(n_{kb} + 1)$ if $i \in \mathcal{I}_{kb}^{(-c^*)}$ and 0 otherwise. Barber Theorem 2 gives

$$P(\text{coverage in cell } (k, b)) \geq 1 - \alpha - \sum_{i \in \mathcal{I}_{kb}^{(-c^*)}} \tilde{w}_i \cdot d_{\text{TV}}(R(Z), R(Z^{(i)})).$$

603 If i is from chromosome c (so $Z_i \sim F_c^{(k,b)}$) and Z_{test} is from c^* (so $Z_{\text{test}} \sim F_{c^*}^{(k,b)}$), the swap
604 $Z_i \leftrightarrow Z_{\text{test}}$ changes the marginal law of exactly two entries of the residual vector. Under 1 (cross-
605 chromosome independence), entries from different chromosomes are independent, so the TV between
606 $R(Z)$ and $R(Z^{(i)})$ is bounded by the TV between the two swapped marginals:

$$d_{\text{TV}}(R(Z), R(Z^{(i)})) \leq d_{\text{TV}}(F_c^{(k,b)}, F_{c^*}^{(k,b)}).$$

607 Substituting and using uniform weights across chromosomes (each of the $|\mathcal{C}| - 1$ calibration chromo-
608 somes contributes equally to the average):

$$P(Y_{\text{test}} \in \mathcal{C}_\alpha(X_{\text{test}}) \mid Y_{\text{test}} = k, b(X_{\text{test}}) = b) \geq 1 - \alpha - \bar{\delta}_{\text{TV}}^{(k,b)},$$

609 where

$$\bar{\delta}_{\text{TV}}^{(k,b)} = \frac{1}{|\mathcal{C}| - 1} \sum_{c \neq c^*} d_{\text{TV}}(F_c^{(k,b)}, F_{c^*}^{(k,b)}).$$

610 The bound has no factor of 2. □

611 **Relation to T3.** Setting $\bar{\delta}_{\text{TV}}^{(k,b)} = 0$ recovers T3's lower bound exactly. T3' degrades T3 by precisely
612 the cell-level chromosome-TV quantity that our cell-level KS audit (App. B.4) empirically bounds.

613 **A.8 Lemma 8: T3-loc feature-ball coverage**

614 **Lemma 8** (T3-loc feature-ball coverage). *Suppose $\hat{\sigma} : \mathcal{X} \rightarrow \mathbb{R}_+$ is L -Lipschitz in the ℓ_2 metric and*
615 *A3-loc (local score stationarity on ℓ_2 -balls) holds. Let $\Delta = \min_b (e_b - e_{b-1})$ be the minimum $\hat{\sigma}$ -bin*
616 *width. For any $x_0 \in \mathcal{X}$, any radius r with $Lr < \Delta/2$, and any class k ,*

$$|P(Y \in \mathcal{C}_\alpha(X) \mid X \in B_\ell(x_0, r), Y = k) - (1 - \alpha)| \leq \frac{1}{n_{\min}(x_0, r) + 1},$$

617 where $n_{\min}(x_0, r)$ is the minimum calibration count over the at-most-two $\hat{\sigma}$ -bins that intersect
618 $\hat{\sigma}(B_\ell(x_0, r))$.

619 *Proof.* Assume $\hat{\sigma} : \mathcal{X} \rightarrow \mathbb{R}_+$ is L -Lipschitz and A3-loc holds. Fix $x_0 \in \mathcal{X}$, $r > 0$ with $Lr < \Delta/2$,
 620 and class k .

621 *Step 1 (Lipschitz bound on $\hat{\sigma}$ -range).* For $x \in B_\ell(x_0, r)$:

$$|\hat{\sigma}(x) - \hat{\sigma}(x_0)| \leq L\|x - x_0\| \leq Lr.$$

622 So $\hat{\sigma}(B_\ell(x_0, r)) \subseteq [\hat{\sigma}(x_0) - Lr, \hat{\sigma}(x_0) + Lr]$.

623 *Step 2 (Pigeonhole — at most two bins).* Since $Lr < \Delta/2$, this $\hat{\sigma}$ -interval has width $2Lr < \Delta$
 624 and therefore intersects at most two adjacent bins $B_{b_\star}, B_{b_\star+1}$. (An interval of width less than the
 625 minimum bin width can straddle at most one bin boundary.)

626 *Step 3 (Convex decomposition).* Decompose the conditional coverage:

$$\begin{aligned} & P(Y \in \mathcal{C}_\alpha(X) \mid X \in B_\ell(x_0, r), Y = k) \\ &= \sum_{b \in \{b_\star, b_\star+1\}} w_b(x_0, r) \cdot P(Y \in \mathcal{C}_\alpha(X) \mid X \in B_\ell(x_0, r), Y = k, b(X) = b), \end{aligned}$$

627 where $w_b(x_0, r) := P(b(X) = b \mid X \in B_\ell(x_0, r), Y = k)$ are non-negative and sum to 1.

628 *Step 4 (Within-cell restriction preserves F_{kb}).* Inside each cell (k, b) , the conformal threshold
 629 $\hat{q}_{kb}$ depends only on calibration and is not affected by restricting X to $B_\ell(x_0, r)$. Under A3-loc,
 630 $s(X, k) \mid (Y = k, b(X) = b, X \in B_\ell(x_0, r))$ has the cell score law F_{kb} . The T3 argument
 631 (App. A.6, Steps 2–4) then applies verbatim inside each cell restricted to the ball:

$$P(Y \in \mathcal{C}_\alpha(X) \mid X \in B_\ell(x_0, r), Y = k, b(X) = b) \in \left[1 - \alpha, 1 - \alpha + \frac{1}{n_{kb} + 1} \right].$$

632 *Step 5 (Convex-hull sandwich).* The convex combination over the (at most two) bins $\{b_\star, b_\star + 1\}$ lies
 633 in

$$\left[1 - \alpha, 1 - \alpha + \frac{1}{n_{\min}(x_0, r) + 1} \right],$$

634 where $n_{\min}(x_0, r) = \min_{b \in \{b_\star, b_\star+1\}} n_{kb}$. □

635 **Honest limitations.** A3-loc is stronger than A2-cell — it requires local stationarity at the scale
 636 r , which we cannot fully verify at high feature-space dimension with our $n \leq 11,400$. We state
 637 T3-loc conditionally on A3-loc. The cell-level KS audit (App. B.4) is the closest finite-sample analog.
 638 The Lipschitz constant L of $\hat{\sigma}$ (distinct from the score-CDF Lipschitz L_F of 4) is not estimated in
 639 this paper and is deferred to future work; T3-loc is therefore stated as a conditional guarantee, not
 640 an operationally-deployed bound. Finally, Barber et al. [2020]’s impossibility is not violated: we
 641 guarantee coverage averaged over $B_\ell(x_0, r)$, not pointwise, and the resolution $r < \Delta/(2L)$ prevents
 642 r from going to zero without re-binning.

643 A.9 Lemma 9: T3.b $\hat{\sigma}$ perturbation

644 **Lemma 9** (T3.b $\hat{\sigma}$ perturbation). *Let $\sigma^\star$ be an oracle variance function and assume B1 ($\|\hat{\sigma} -$
 645 $\sigma^\star\|_\infty / \inf_x \sigma^\star(x) \leq \eta$). Under A1' + A2-cell (w.r.t. the fixed bin edges $e_0 < \dots < e_K$),*

$$\left| P(Y \in \mathcal{C}_\alpha^{(\hat{\sigma})}(X) \mid Y = k) - P(Y \in \mathcal{C}_\alpha^{(\sigma^\star)}(X) \mid Y = k) \right| \leq \frac{2\eta \bar{\sigma}}{\Delta} + \frac{1}{n_{\min} + 1},$$

646 where $\bar{\sigma} = \sup_x \sigma^\star(x)$ and $n_{\min}$ is the minimum cell size.

647 *Proof.* Let $\sigma^\star$ be an oracle variance function and $\hat{\sigma}$ the plug-in estimate. The two predictors $\mathcal{C}_\alpha^{(\hat{\sigma})}$ and
 648 $\mathcal{C}_\alpha^{(\sigma^\star)}$ differ in two ways:

649 (a) *Bin reassignment.* A variant x with $\sigma^\star(x) = e_b + t$ for small $t > 0$ (near a bin boundary)
 650 may be placed in bin b under $\sigma^\star$ and bin $b + 1$ under $\hat{\sigma}$, whenever $|\hat{\sigma}(x) - \sigma^\star(x)| > t$. Under B1,
 651 $|\hat{\sigma}(x) - \sigma^\star(x)| \leq \eta \cdot \sigma^\star(x) \leq \eta \bar{\sigma}$ where $\bar{\sigma} = \sup_x \sigma^\star(x)$. A variant is at risk of bin reassignment iff
 652 $|\sigma^\star(x) - e_b| < \eta \bar{\sigma}$ for some bin edge e_b . By Lemma 4 (with $\delta = \eta \bar{\sigma}$), the total probability mass of
 653 such variants is bounded by $2\eta \bar{\sigma} / \Delta$.

654 (b) *Score rescaling*. Within the same cell, the scores differ:

$$s_{\hat{\sigma}}(x, k) = \frac{|k - \hat{p}(x)|}{\hat{\sigma}(x) + \epsilon} = s_{\sigma^*}(x, k) \cdot \frac{\sigma^*(x) + \epsilon}{\hat{\sigma}(x) + \epsilon} \in [(1 - \eta) s_{\sigma^*}(x, k), (1 + \eta) s_{\sigma^*}(x, k)].$$

655 The rescaling function $g : s \mapsto s \cdot (\sigma^*(x) + \epsilon) / (\hat{\sigma}(x) + \epsilon)$ is strictly increasing. By Lemma 3, both
 656 the calibration quantile and the test score transform together under the same monotone map, so the
 657 coverage indicator is invariant. Contribution (b) is exactly zero.

658 (c) *Combining*. The coverage difference between the plug-in and oracle predictors is bounded by (a)
 659 plus the T3 finite-sample slack within unaffected cells:

$$\left| P(Y \in \mathcal{C}_{\alpha}^{(\hat{\sigma})}(X) \mid Y = k) - P(Y \in \mathcal{C}_{\alpha}^{(\sigma^*)}(X) \mid Y = k) \right| \leq \frac{2\eta \bar{\sigma}}{\Delta} + \frac{1}{n_{\min} + 1}. \quad \square$$

660 **Graceful degradation.** As $\hat{\sigma} \rightarrow \sigma^*$ uniformly ($\eta \rightarrow 0$), the T3.b bound collapses to the T3
 661 finite-sample slack $1/(n_{\min} + 1)$. The plug-in conformal predictor is thus continuous in the quality
 662 of $\hat{\sigma}$ — there is no cliff where a slightly worse $\hat{\sigma}$ breaks coverage.

663 **A.10 Lemma 10: T4 chrom-shift bound (Barber 2023 Thm 2)**

664 **Lemma 10** (T4 chrom-shift bound, Barber et al. [2023] Thm 2). *For any prespecified weights*
 665 *$w_i \in [0, 1]$ and $\tilde{w}_i = w_i / (w_1 + \dots + w_n + 1)$,*

$$P(Y \in \mathcal{C}_{\alpha}(X)) \geq 1 - \alpha - \sum_{i=1}^n \tilde{w}_i \cdot d_{\text{TV}}(R(Z), R(Z^{(i)})),$$

666 *where $R(Z) = (S_1, \dots, S_{n+1})$ is the score vector and $R(Z^{(i)})$ is the sequence with Z_i and Z_{n+1}*
 667 *swapped.*

668 *Proof.* This is a direct restatement of Barber et al. [2023] Theorem 2 in our notation.

669 Let $Z = (Z_1, \dots, Z_n, Z_{n+1})$ with $Z_{n+1} = (X_{\text{test}}, Y_{\text{test}})$ the test point, and $R(Z) = (S_1, \dots, S_{n+1})$
 670 the score vector. For any prespecified weights $w_i \in [0, 1]$ with $\tilde{w}_i = w_i / (w_1 + \dots + w_n + 1)$ and
 671 conformal predictor $\hat{C}_n$ built from a weighted quantile of $\{S_i\}_{i=1}^n$, Barber et al. [2023] show:

$$P(Y_{n+1} \in \hat{C}_n(X_{n+1})) \geq 1 - \alpha - \sum_{i=1}^n \tilde{w}_i \cdot d_{\text{TV}}(R(Z), R(Z^{(i)})),$$

672 *where $R(Z^{(i)})$ is the score vector with Z_i and Z_{n+1} swapped.*

673 Under 1, entries from different chromosomes are independent, and the swap $Z_i \leftrightarrow Z_{n+1}$ changes the
 674 marginal law of exactly two entries. If Z_i comes from chromosome c and Z_{n+1} from c^* :

$$d_{\text{TV}}(R(Z), R(Z^{(i)})) \leq d_{\text{TV}}(F_c, F_{c^*}).$$

675 The robust corollaries T1', T2', T3' are direct consequences of this bound applied at the appropriate
 676 partition level:

- 677 • T1': uniform weights, marginal scores, yields $\bar{\delta}_{\text{TV}}$ (App. A.5);
- 678 • T2': restrict to class- k calibration indices, class-conditional scores, yields $\bar{\delta}_{\text{TV}}^{(k)}$;
- 679 • T3': restrict to cell- (k, b) calibration indices, cell-conditional scores, yields $\bar{\delta}_{\text{TV}}^{(k, b)}$ (App. A.7).

680 The bound carries no factor of 2. This is the weighted-swap construction of Barber et al. [2023], not
 681 the jackknife+ bound from an earlier Barber et al. [2021] paper. $\square$

682 **A.11 Proof of Theorem 3 (T5.1 oracle bin count)**

683 We first establish the bias–variance decomposition, then derive the oracle K^* .

684 *Proof of the gap decomposition (Eq. (3)).* Under 2 and 4:

685 *Step 1 (Bias bound).* Within bin b , all $\hat{\sigma}$ values lie in $(e_{b-1}, e_b]$ with width Δ_b . The cell score CDF
686 F_{kb} is a mixture:

$$F_{kb}(s) = \int_{e_{b-1}}^{e_b} F_{k,\sigma}(s) dG_b(\sigma),$$

687 where G_b is the distribution of $\hat{\sigma}$ within bin b . By 4, for any reference value $\sigma_0 \in (e_{b-1}, e_b]$:

$$\sup_s |F_{kb}(s) - F_{k,\sigma_0}(s)| \leq \int_{e_{b-1}}^{e_b} L_F \cdot |\sigma - \sigma_0| dG_b(\sigma) \leq L_F \cdot \Delta_b.$$

688 Now consider two calibration chromosomes c, c' with score laws $F_c^{(k,b)}$ and $F_{c'}^{(k,b)}$. Under 2, within
689 each chromosome, the score law conditional on cell (k, b) is the chromosome-specific conditional
690 CDF. If 4 holds uniformly across chromosomes (i.e., L_F does not depend on c), both $F_c^{(k,b)}$ and
691 $F_{c'}^{(k,b)}$ lie within KS-distance $L_F \cdot \Delta_b$ of the same reference law F_{k,σ_0} . By T3' (App. A.7), the
692 coverage deviation due to within-bin heterogeneity is bounded by $L_F \cdot \Delta_b$. This is the stationarity
693 bias $\beta_{kb} = L_F \cdot \Delta_b$.

694 *Step 2 (Variance bound).* Under exact A2-cell (or using T3' with TV = 0), T3 gives the finite-sample
695 gap $\leq 1/(n_{kb} + 1)$. This is the variance term $\nu_{kb} = 1/(n_{kb} + 1)$.

696 *Step 3 (Combining).* When A2-cell holds only approximately with bias β_{kb} , the total gap is at most
697 $\beta_{kb} + \nu_{kb}$.

698 For equiprobable bins, $\Delta_b = R/K$ and $n_{kb} \geq \lfloor \pi_{\min} \cdot n/K \rfloor$, so the worst-cell gap satisfies

$$G(K) \leq \frac{L_F \cdot R}{K} + \frac{K}{\pi_{\min} \cdot n}.$$

699 (The floor is absorbed into the leading $O(n^{-1/2})$ term for the oracle.) $\square$

700 *Proof of Theorem 3 (T5.1).* Treat K as a continuous variable. The gap function $g(K) = L_F R/K +$
701 $K/(\pi_{\min} n)$ is convex, so setting $g'(K) = 0$:

$$-\frac{L_F R}{K^2} + \frac{1}{\pi_{\min} n} = 0 \implies K^* = \sqrt{L_F \cdot R \cdot \pi_{\min} \cdot n}.$$

702 Substituting back:

$$g(K^*) = \frac{L_F R}{\sqrt{L_F R \cdot \pi_{\min} n}} + \frac{\sqrt{L_F R \cdot \pi_{\min} n}}{\pi_{\min} n} = 2\sqrt{\frac{L_F R}{\pi_{\min} n}}.$$

703 The floor $\lfloor K^* \rfloor$ introduces an additive $O(1/(\pi_{\min} n))$ term (the gap between K^* and $\lfloor K^* \rfloor$ shifts the
704 variance term by at most $1/(\pi_{\min} n)$), which is dominated by the $O(n^{-1/2})$ rate. $\square$

705 **A.12 Proof of Theorem 4 (T5.2 matching lower bound)**

706 *Proof.* We construct a distribution achieving both lower bounds simultaneously.

707 *Bias lower bound.* Consider the family $F_{k,\sigma}(s) = \Phi(s - L_F \sigma)$, where Φ is the standard normal CDF.
708 This satisfies 4 with Lipschitz constant L_F : for any σ_1, σ_2 ,

$$\sup_s |F_{k,\sigma_1}(s) - F_{k,\sigma_2}(s)| = \sup_s |\Phi(s - L_F \sigma_1) - \Phi(s - L_F \sigma_2)| \leq L_F |\sigma_1 - \sigma_2|,$$

709 where the inequality uses $\|\Phi'\|_\infty = \phi(0) = 1/\sqrt{2\pi} < 1$.

710 In any bin of width $\Delta_b = R/K$, the mixture CDF F_{kb} averages over shifted CDFs. Because the
711 shift is monotone in σ , the CDF at the lower boundary $F_{k,e_{b-1}}$ and upper boundary F_{k,e_b} straddle the
712 mixture:

$$\sup_s |F_{kb}(s) - F_{k,e_{b-1}}(s)| \geq \frac{L_F \Delta_b}{2} = \frac{L_F R}{2K}.$$

(The factor $1/2$ comes from the average shift relative to one endpoint being at least half the bin width.)

Variance lower bound. For any cell with n_{kb} calibration points, the conformal quantile’s rank-based coverage is exactly $\lceil (n_{kb} + 1)(1 - \alpha) \rceil / (n_{kb} + 1)$, which lies in $[1 - \alpha, 1 - \alpha + 1/(n_{kb} + 1)]$. When F_{kb} is continuous and the test point’s rank is uniformly distributed, the gap $1/(n_{kb} + 1)$ is attainable (up to ceiling effects). For the worst cell, $n_{kb} \leq \lceil \pi_{\min} n / K \rceil$, so the gap is at least $1/(2\lceil \pi_{\min} n / K \rceil + 1)$.

Combined lower bound. For any K :

$$G(K) \geq \max\left(\frac{L_F R}{2K}, \frac{1}{2\lceil \pi_{\min} n / K \rceil + 1}\right).$$

Optimizing over K : the first term is $\Theta(1/K)$ and the second is $\Theta(K/n)$. They balance at $K = \Theta(\sqrt{n})$, giving

$$\min_K G(K) = \Omega\left(\sqrt{\frac{L_F R}{\pi_{\min} n}}\right),$$

matching the $O(n^{-1/2})$ upper bound of T5.1 up to constants. $\square$

Interpretation. The $O(n^{-1/2})$ rate is dimension-free because T5 conditions on the 1-D summary $\hat{\sigma}(x)$, not on the full feature vector. By contrast, RLCP [Hore and Barber, 2025] achieves local coverage at rate $O(n^{-2/(d+2)})$, which is impractical at $d = 7,731$. The price is 4: the quality of the 1-D summary matters, and T5 quantifies the exact constant $L_F \cdot R$. In the homoscedastic limit $L_F = 0$, the oracle $K^* = 0$ and no binning is needed — T1/T2 suffice.

B Data and training details

B.1 TraitGym benchmark

We use TraitGym [Benegas et al., 2025], a standardized evaluation suite for variant effect prediction. Two matched datasets:

- **Mendelian (matched_9):** 338 ClinVar pathogenic variants + 3,042 matched benign controls = 3,380 total across 113 OMIM phenotypes. Matching on consequence type, minor allele frequency (MAF), and transcription start site (TSS) distance (9 strata).
- **Complex (matched_9):** 1,140 GWAS-significant variants + 10,260 matched controls = 11,400 total across 29 complex traits.

Chrom-LOO protocol. Test chromosomes: $\{17, 18, 19, 20, 21, 22, X\}$ (TraitGym default). For each test chrom c^* , the remaining chromosomes form the train/calibration fold. Primary metric: $\text{AUPRC}_{\text{chrom-weighted}} = \sum_c w_c \cdot \text{AUPRC}_c$ with $w_c = n_c / \sum_{c'} n_{c'}$.

B.2 Input features

B.3 Aggregator hyperparameters

Base classifier $\hat{p}$. HistGradientBoostingClassifier (scikit-learn 1.4): `max_depth=2`, `max_iter=100`, `learning_rate=0.1`, `class_weight=balanced`, `random_state=42`, `early_stopping=False`. Feature imputation: mean via SimpleImputer. No feature scaling (GBM is scale-invariant).

AlphaGenome features. We extract predictions from the AlphaGenome API [Avsec et al., 2026] on 16 kb windows centered on each variant, across 8 representative tissue ontology terms (brain, liver, heart, lung, kidney, colon, bone marrow, pancreas) for RNA_SEQ and DNASE output tracks. For each (output, tissue) pair we record (a) window-mean $\log(1 + x)$ of ref and alt predictions, (b) center-position alt-ref delta in log space, and (c) maximum $|\text{alt} - \text{ref}|$ across the window, yielding $2 \cdot 8 \cdot (2 + 2) = 64$ entries per output type, 128 combined — plus 12 aggregated cross-tissue summaries

Table 3: Features used in HCCP aggregator. All pre-computed as per TraitGym conventions.

Feature	Source	Dim
CADD v1.7	Kircher et al. [2014]	1
Borzoï_L2	Linder et al. [2025] zero-shot L2-norm	3,846
Borzoï_L2_L2	Linder et al. [2025] L2-of-L2	3,846
GPN-MSA_LLRL	Benegas et al. [2023] log-likelihood ratio	1
GPN-MSA_absLLR	Benegas et al. [2023] absolute LLR	1
GPN-MSA_InnerProducts	Benegas et al. [2023] inner products	36
AlphaGenome	Avsec et al. [2026] RNA_SEQ + DNASE $\times$ 8 tissues	140
CADD+GPN-MSA+Borzoï (full)		7,731
CADD+GPN-MSA+Borzoï+AlphaGenome		7,871
CADD+Borzoï		7,693

for 140 total. Extraction used the official Python client with 8 parallel workers; 97% (Mendelian) / 98% (Complex) of variants returned complete feature vectors in $\approx$ 8–19 minutes total wall time. Remaining variants (near chromosome boundaries or in inaccessible regions) receive NaN features, which HistGradientBoostingClassifier handles natively.

Heteroscedastic head $\hat{\sigma}$. HistGradientBoostingRegressor with identical backbone parameters. Target: $r_i = |Y_i - \hat{p}^{(-C_i)}(X_i)|$ (chrom-LOO OOF residuals). Loss: squared error (Gaussian NLL approximation). Floor: $\hat{\sigma}(x) = \max(\hat{\sigma}_{\text{raw}}(x), \epsilon)$ with $\epsilon = 10^{-4}$.

Conformal calibration. Score: $s(x, y) = |y - \hat{p}(x)|/(\hat{\sigma}(x) + \epsilon)$. Bin count: $K = \hat{K}(c_{\text{outer}})$ via nested chrom-LOO (§5.2); fallback $K = 5$ for fixed-K baselines (Tab. 2). Cell minimum $n_{k,b} \geq 5$; fallback to class-pooled $\hat{q}_k$ otherwise. Miscoverage: $\alpha = 0.10$.

DEGU-lite ensemble. $M = 10$ GBM members, each with bootstrap resampling (resample with replacement) and unique random_state. Ensemble uncertainty: $\hat{\sigma}_{\text{DEGU}}(x) = \text{std}_{m=1}^M \hat{p}_m(x)$. Same backbone as aggregator.

B.4 Within-cell KS audit

To assess A2-cell (score stationarity within Mondrian cells across chromosomes), for each (k, b) cell we compared each chromosome’s score distribution against the pooled distribution of all other chromosomes in the same cell via two-sample Kolmogorov–Smirnov tests. Cells with $n \geq 20$ per-chrom were retained (52 tests on Mendelian, 125 on Complex, matching the $\approx$ 155-test budget reported in earlier drafts aggregated across datasets).

Table 4: Per-chromosome A2-cell KS audit (CADD+GPN-MSA+Borzoï). A perfectly stationary score law would reject at the nominal 5% rate. The score law $s(x, y) = |y - \hat{p}(x)|/(\hat{\sigma}(x) + \epsilon)$ is used.

Dataset	n_{tests}	Reject @ $p < 0.05$	KS max	KS mean	KS median
Mendelian	52	46.2% (24/52)	0.434	0.207	0.163
Complex	125	10.4% (13/125)	0.315	0.113	0.102

Interpretation. A2-cell is approximately satisfied on Complex (10.4% vs 5% null, max KS 0.315) but noticeably violated on Mendelian (46.2%, max KS 0.434). The violation concentrates in the highest- $\hat{\sigma}$ positive-class cell ($y=1, b=4$) — chromosomes chr9, chr7, chr3, chr19, and chrX each show $\text{KS} \geq 0.40$ at this cell, consistent with chromosome-specific pathogenic variant concentration. The robust form T3’ absorbs this violation via the chrom-TV term: our observed per-chrom max marginal-coverage gap of 0.100 (chr6, Tab. 16) is within the $\bar{\delta}_{\text{TV}}^{(k,b)}$ bound consistent with $\max \text{KS} \leq 0.44 \Rightarrow d_{\text{TV}} \leq 0.22$. Full per-chromosome breakdowns: outputs/ks_audit/*_ks_per_chrom_n20.json.

780 ***K*-invariance of the rejection rate.** To confirm that the 46.2% rejection on Mendelian is a
781 dataset-intrinsic chrom-TV signal rather than a binning artefact at the legacy default $K = 5$, we
782 re-run the per-chromosome KS audit on Mendelian at $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$ with the same min-
783 per-chrom-cell threshold of 20 (Tab. 5). The rejection rate stays in the 40–46% band across all
784 five K values, and the maximum cell-by-chrom KS statistic stays in $[0.434, 0.565]$. T3' is there-
785 fore the binding operational guarantee on Mendelian regardless of partition granularity. Source:
786 `outputs/ks_audit_kcv/mendelian_K{2,3,5,8,10}_n20.json`.

Table 5: Mendelian per-chromosome A2-cell KS audit K -sweep at min-per-chrom-cell= 20. Rejec-
tion rate is the fraction of (cell, chromosome) pairs at which $p < 0.05$.

K_{eval}	n_{tests}	Reject @ $p < 0.05$	KS max
2	37	45.9% (17/37)	0.565
3	45	40.0% (18/45)	0.524
5	52	46.2% (24/52)	0.434
8	58	43.1% (25/58)	0.492
10	60	45.0% (27/60)	0.558

787 B.5 Leakage sanity audit

788 Because HistGradientBoostingClassifier at max_depth=2 with early_stopping=False,
789 default subsample=1.0, and max_features=None contains no active random source, the classifier
790 output is *deterministic given the training data* and random_state produces no variance. We view
791 this as an algorithmic property (sklearn 1.4 documentation confirms that random_state controls
792 only early-stopping validation splits, row subsampling, and binning tie-breaking, none of which
793 activate under our configuration), but algorithmic determinism is observationally indistinguishable
794 from *data leakage* absent additional tests. We therefore report three orthogonal sanity checks on
795 CADD+GPN-MSA+Borzo Mendelian and CADD+Borzo Complex.

Table 6: Leakage sanity audit. Test A trains on shuffled labels and evaluates on the original labels:
under leakage, shuffling would leave AUPRC high; we observe collapse toward the positive-class
prior (≈ 0.10). Test B forces row-level stochasticity via 80% random row masking with seeds
 $\{42, 7, 2024\}$ and confirms random_state *does* produce variance under active randomness — but
less than 0.003 AUPRC std, reflecting TraitGym’s strong feature signal rather than a bug. Test C
audits the chrom-LOO split.

Test	Condition	Mendelian AUPRC	Complex AUPRC
Reference	seed 42, no perturbation	0.900	0.350
A: label shuffle	10% labels flipped	0.898	0.327
A: label shuffle	30% labels flipped	0.776	0.257
A: label shuffle	50% labels flipped	0.147	0.126
B: subsample+seed	row-mask 0.8 (std over 3 seeds)	0.0006 (std)	0.0027 (std)
C: chrom-LOO audit	train $\cap$ test, folds 19–22	0 overlap	0 overlap

796 **Interpretation.** (i) *Test A is the critical test:* on both datasets, training with 50% random label
797 corruption collapses test AUPRC from 0.90/0.35 to 0.147/0.126, approaching the positive-class base
798 rate 0.10. This is inconsistent with data leakage — a leaked model would retain its AUPRC regardless
799 of training labels — and consistent with the model genuinely learning label signal. (ii) Test B shows
800 that when stochasticity is forcibly introduced via row subsampling, seed variance is strictly positive
801 (0.0006 on Mendelian, 0.0027 on Complex). The asymmetry is biologically interpretable: Complex
802 has weaker signal (AUPRC 0.35 vs 0.90), so the model is more sensitive to training perturbations,
803 as expected. The fact that both stds are small (< 0.003) reflects TraitGym’s strong pre-computed
804 features (CADD, Borzoi, GPN-MSA) — the majority of predictive signal comes from the features
805 themselves, not the boosting trajectory. (iii) Test C confirms the chrom-LOO split is structurally
806 clean. Together, the three tests establish that the zero seed-variance observation in the main text is an
807 algorithmic property, not a leakage artifact, and that chromosome-level bootstrap (App. B.6) is the
808 appropriate uncertainty quantifier.

809 B.6 Compute

810 All experiments on a single machine with $4 \times$ NVIDIA T4 GPUs (SM 7.5, no BF16). The HCCP
811 pipeline (aggregator + $\hat{\sigma}$ head + conformal calibration) runs entirely on CPU via scikit-learn.

Table 7: Wall-clock cost of the full HCCP pipeline and main baselines, single chrom-LOO sweep on TraitGym. CPU times on 16-core Xeon; GPU on a single T4. Memory peak excludes Python overhead. $\hat{p}$ training is per feature set; $\hat{\sigma}$ training is on chrom-LOO OOF residuals; Mondrian calibration is reported per K . DEGU-lite shares the $\hat{p}$ backbone. The $B = 200$ chromosome-level bootstrap (Tab. 2) reuses cached calibration scores and adds ~ 30 CPU-min total.

Stage	Backend	Mendelian	Complex	Mem peak
$\hat{p}$ training (HistGB)	CPU 16-core	≈ 20 min	≈ 60 min	4 GB
$\hat{\sigma}$ training (HistGB Gaussian-NLL)	CPU 16-core	≈ 15 min	≈ 45 min	4 GB
Mondrian conformal calibration	CPU single-thread	< 5 min	< 5 min	1 GB
Total HCCP (single feature set, all chroms)	CPU 16-core	< 4 CPU-h	< 4 CPU-h	4 GB
DEGU-lite ($M = 10$ bootstrap GBM)	CPU 16-core	≈ 35 min	≈ 75 min	6 GB
$B = 200$ chrom-bootstrap	CPU single-thread	≈ 12 min	≈ 18 min	1 GB
Multi-seed sweep ($\{42, 7, 2024\}$)	CPU 16-core	≈ 12 min	≈ 18 min	4 GB

812 Figure and proof generation are negligible (matplotlib + LaTeX).

813 C Per-partition tables

814 C.1 Subsumption table: HCCP vs. contemporary / closest CP work

815 Tab. 8 crystallises the prose comparison in §2 into a single per-axis grid. Each row marks whether
816 the method explicitly addresses (a) a learned heteroscedastic head $\hat{\sigma}(x)$ with finite-sample local
817 guarantee, (b) class-conditional coverage, (c) coverage local on the $\hat{\sigma}$ axis (rather than a different
818 feature axis or kernel), (d) an $\pi_{\min}$ -explicit term in the worst-cell rate, (e) a matching within-class
819 lower bound, and (f) a data-dependent partition (i.e., the partition itself is computed from data rather
820 than pre-specified). HCCP is the only row marked $\checkmark$ on all six axes. The Yao 2026 / Bairaktari 2025
821 / Plassier 2025 concurrent works each cover a strict subset.

Table 8: HCCP vs. closest / concurrent CP work, by guarantee axis. $\checkmark$ = method explicitly addresses the axis with a finite-sample guarantee; $\circ$ = handled but without a finite-sample bound, or on a different axis (e.g., RLCP: kernel localisation in feature space, not the $\hat{\sigma}$ axis); $\times$ = not addressed. “Lower bound” refers to a matching within-class minimax lower bound that pairs with the upper-bound rate; Yao 2026 / xie 2025 do not provide one for their respective rates. “Data-dep. partition” refers to a Mondrian-style partition computed from data; class-conditional CP [Sadinle et al., 2019] and Boström 2020 use pre-specified groups (label / $\hat{\sigma}$ -quantile bins fixed before calibration). Bairaktari et al. [2025] subsume class-conditional and Mondrian as static-group special cases but require pre-specified groups, hence $\times$ on the partition column. SC-CP [van der Laan, 2024] uses an isotonic recalibration on the $\hat{\mu}$ axis (different from $\hat{\sigma}$).

Method	Hetero. $\hat{\sigma}$	Class-cond.	Local- $\hat{\sigma}$	$\pi_{\min}$ in rate	Lower bd.	Data-dep. part.
Sadinle et al. [2019] (LAC)	$\times$	$\checkmark$	$\times$	$\times$	$\times$	$\times$
Boström and Johansson [2020]	$\checkmark$	$\times$	$\checkmark$	$\times$	$\times$	$\times$
Hore and Barber [2025] (RLCP)	$\circ$	$\times$	$\circ$ kernel $\mathbf{x}$	$\times$	$\times$	$\times$
van der Laan [2024] (SC-CP)	$\circ$	$\times$	$\circ$ $\hat{\mu}$ -axis	$\times$	$\times$	$\checkmark$ (isotonic)
Xie et al. [2025]	$\times$	$\checkmark$	$\times$	$\circ$ marginal	$\times$	$\times$
Bairaktari et al. [2025] (Kandinsky)	$\times$	$\checkmark$	$\circ$ overlapping	$\times$	$\times$	$\times$
Yao et al. [2026]	$\times$	$\times$	$\times$	$\times$	$\times$	$\times$
Plassier et al. [2025] (CP ² -HPD)	$\circ$	$\times$	TV-density-ratio	$\times$	$\times$	$\times$
HCCP (ours)	$\checkmark$	$\checkmark$	$\checkmark$ $\hat{\sigma}$ -bin	$\checkmark$ T5.1	$\checkmark$ T5.2	$\checkmark$ ($y \times \hat{\sigma}$)

822 C.2 Single-axis CP head-to-head and contemporary CP comparison (deferred from §6.2)

823 Tab. 9 reports the four single-axis Mondrian variants (split CP, $\hat{\sigma}$ -Mondrian, class-Mondrian, HCCP)
824 on TraitGym; Tab. 10 compares HCCP against contemporary local-CP methods (RLCP, weighted CP,
825 SC-CP) with $B = 200$ chrom-level bootstrap CIs. Tab. 11 reports the imbalance sweep on Complex.

Table 9: Head-to-head against off-the-shelf normalized + Mondrian CP, sharing the same $(\hat{p}, \hat{\sigma})$. Baselines: B2 [Boström and Johansson, 2020], B3 [Sadinle et al., 2019]. Partition K for B2 / B4 selected per outer fold via proper nested chrom-LOO with K_{eval} -fair scoring (Mendelian $K_{\text{eval}} = 3$, Complex $K_{\text{eval}} = 5$). **All four rows on $B = 200$ chromosome-level bootstrap (mean $\pm$ std).** The reported $\hat{\sigma}$ -bin gap is the cell-level worst- (k, b) deviation from 0.90, the same object as $G(K)$ in Eq. (3); this differs from the marginal-bin metric of Tab. 2. Of the four variants we evaluated, HCCP is the only frontier point at this corner on both datasets. Note that the on-axis baseline B2 $\hat{\sigma}$ -Mondrian still loses on $\hat{\sigma}$ -bin gap because $\hat{\sigma}$ -only stratification does not preserve $\text{cov}|_{Y=1}$ at $\pi_{\min} = 0.10$; the $(y \times \hat{\sigma})$ product is necessary, not just the $\hat{\sigma}$ axis.

Dataset	Method	Marg. cov.	$\text{Cov} _{Y=1}$	$\hat{\sigma}$ -bin gap $\downarrow$
Mendelian ($\pi_+ = 0.10$)	B1 vanilla split CP	0.899 $\pm$.009	0.982$\pm$.008	0.265 $\pm$.153
	B2 $\hat{\sigma}$ -Mondrian	0.901 $\pm$.013	0.821$\pm$.049	0.401$\pm$.187
	B3 class-Mondrian	0.898 $\pm$.008	0.895 $\pm$.032	0.410$\pm$.172
	B4 HCCP (ours)	0.900 $\pm$.012	0.875$\pm$.040	0.173$\pm$.192
Complex ($\pi_+ = 0.10$)	B1 vanilla split CP	0.900 $\pm$.003	0.614$\pm$.018	0.822$\pm$.028
	B2 $\hat{\sigma}$ -Mondrian	0.900 $\pm$.003	0.624$\pm$.018	0.827$\pm$.025
	B3 class-Mondrian	0.900 $\pm$.004	0.900 $\pm$.016	0.805$\pm$.051
	B4 HCCP (ours)	0.900 $\pm$.003	0.905$\pm$.010	0.060$\pm$.031

Table 10: HCCP vs. contemporary local-CP methods, sharing $(\hat{p}, \hat{\sigma})$. Baselines: RLCP [Hore and Barber, 2025], weighted CP [Tibshirani et al., 2019], SC-CP [van der Laan, 2024]. $\hat{\sigma}$ -bin gap is the cell-level worst- (k, b) metric (matched to Tab. 9 convention). Subscripts are $\pm$ one std from $B = 200$ chrom-level bootstrap resamples. Complex CI: HCCP [0.019, 0.119] vs. baselines [0.770, 0.879] (disjoint).

Dataset	Method	Marg. cov.	$\text{Cov} _{Y=1}$	$\hat{\sigma}$ -bin gap $\downarrow$	Singleton
Complex ($K_{\text{eval}} = 5$)	HCCP (ours)	0.900	0.905$\pm$.010	0.060$\pm$.031	0.224
	RLCP	0.952	0.543$\pm$.019	0.820$\pm$.029	0.762
	weighted CP	0.901	0.611$\pm$.020	0.821$\pm$.030	0.785
	SC-CP	0.900	0.625$\pm$.018	0.817$\pm$.034	0.710
Mendelian ($K_{\text{eval}} = 3$)	HCCP (ours)	0.902	0.875$\pm$.040	0.173$\pm$.192	0.827
	RLCP	0.939	0.878 $\pm$.034	0.360$\pm$.171	0.890
	weighted CP	0.902	0.982$\pm$.006	0.252 $\pm$.153	0.499
	SC-CP	0.900	0.859$\pm$.044	0.404$\pm$.182	0.777

Table 11: Imbalance sweep: $\text{cov}|_{Y=1}$ as a function of π_+ , Complex CADD+GPN-MSA+Borzoï. Partition K for B2 / B4 selected per outer fold via nested chrom-LOO at $K_{\text{eval}} = 5$. B1 and B2 collapse as $\pi_+ \rightarrow 0$; B3 and B4 are insensitive.

π_+	0.05	0.10	0.20	0.30	0.50	B4 $\hat{\sigma}$ -gap
n	10800	11400	5700	3800	2280	
B1 (split)	0.587	0.616	0.664	0.725	0.810	—
B2 ($\hat{\sigma}$ -Mond)	0.650	0.656	0.718	0.800	0.878	—
B3 (class-Mond)	0.904	0.902	0.902	0.902	0.902	—
B4 (HCCP)	0.898	0.905	0.894	0.902	0.896	0.036 / 0.024 / 0.067 / 0.057 / 0.041

826 C.3 Paired bootstrap test (deferred from §6.2)

827 The marginal $\pm 1\sigma$ bands in Tab. 10 are computed independently per method, so two methods' bands
828 can overlap even when their per-resample diff is consistently negative (the noise is correlated across

829 methods through the shared chrom resample). We therefore re-run the bootstrap with a single seeded
830 RNG drawing chrom indices once per replicate, computing the $\hat{\sigma}$ -bin gap for HCCP and each baseline
831 on the *same* subsample, and recording the paired diff $\Delta_b = \text{gap}_{\text{HCCP},b} - \text{gap}_{\text{baseline},b}$.

Table 12: Paired bootstrap: HCCP $\hat{\sigma}$ -bin gap minus baseline $\hat{\sigma}$ -bin gap, computed on the same $B = 200$ chrom-resamples per replicate (seed 42, identical to Tab. 10). Negative paired mean $\Rightarrow$ HCCP narrower; “Win %” is the share of resamples on which HCCP attained a strictly smaller gap; p_{one} is the one-sided bootstrap p -value (floor $1/(B + 1) \approx 0.005$). Complex shows 100% paired wins at $p \approx 0.005$; on Mendelian, HCCP wins the majority of resamples but the paired CI crosses zero, so we cannot reject equivalence with weighted CP at the 95% level.

Dataset	Comparison	Paired mean	Paired std	CI low	CI high	Win %	p_{one}
Complex ($K_{\text{eval}} = 5$)	HCCP – RLCP	−0.760	0.043	−0.845	−0.678	100.0%	0.005
	HCCP – Weighted CP	−0.760	0.043	−0.845	−0.678	100.0%	0.005
	HCCP – SC-CP	−0.756	0.047	−0.845	−0.663	100.0%	0.005
Mendelian ($K_{\text{eval}} = 3$)	HCCP – RLCP	−0.187	0.141	−0.450	0.000	82.0%	0.184
	HCCP – Weighted CP	−0.078	0.073	−0.191	+0.065	74.5%	0.259
	HCCP – SC-CP	−0.231	0.168	−0.556	0.000	84.5%	0.159

832 **Interpretation.** On Complex, every paired comparison reaches the bootstrap floor ($p \approx 0.005$,
833 100% wins, paired CIs disjoint from zero), confirming that the marginal-CI disjoint claim of Tab. 10
834 understates the strength of the H2H signal. On Mendelian, the paired std (e.g. 0.073 for HCCP-
835 vs-wCP) is roughly one-third of the marginal std on either method individually (0.192 and 0.153),
836 so the marginal-CI overlap was largely a marginal-noise artefact; nevertheless, the paired CI still
837 crosses zero on all three Mendelian comparisons. Operationally we interpret the Mendelian result as:
838 HCCP *tilts* toward narrower gaps on the majority of resamples (74–85%) but does not formally reject
839 equivalence at $\alpha = 0.05$ within $B = 200$. This is consistent with the regime-boundary argument
840 of §6.2: at the per-cell minority floor on Mendelian ($n_{k,b} \approx 100$), the variance term in Eq. (3) is
841 large and erodes finite-sample power for paired tests. The Open Targets cross-platform replication
842 (App. G) reproduces this regime structure: 200/200 paired wins ($p = 0.005$) at $K_{\text{eval}} = 2$ in the
843 deeply bias-dominated regime, 186/200 ($p = 0.075$) at the recommended $K_{\text{eval}} = 5$, and 16/200
844 ($p = 0.92$) at $K_{\text{eval}} = 10$ where HCCP loses (per-cell minority ≈ 114).

845 C.4 Threshold sensitivity for the “Pareto-uniqueness” claim (deferred from §6.2)

846 §6.2 claims HCCP is the only single-fold partition we evaluated that simultaneously meets $\text{cov}_{|Y=1} \geq$
847 0.85 and $\hat{\sigma}$ -bin gap ≤ 0.20 at the recommended K_{eval} . The thresholds (0.85, 0.20) are not unique; we
848 evaluate uniqueness across a 3×4 grid of (T_1, T_2) where $T_1 \in \{0.80, 0.85, 0.90\}$ (minority-coverage
849 floor) and $T_2 \in \{0.15, 0.20, 0.25, 0.30\}$ (cell-level gap ceiling), using the Tab. 9 and Tab. 10 bootstrap
850 means across 7 methods (B1 split, B2 $\hat{\sigma}$ -Mond, B3 class-Mond, HCCP, RLCP, wCP, SC-CP).

Table 13: Methods that meet ($\text{cov}_{|Y=1} \geq T_1$, $\hat{\sigma}$ -bin gap $\leq T_2$) at the recommended K_{eval} , across a $T_1 \times T_2$ grid. “H” = HCCP only; “H, X, ...” = HCCP plus other methods also pass; “—” = no method passes (the ceiling is below HCCP’s own performance). Complex: HCCP is uniquely Pareto-frontier across the entire reasonable grid (gap separation 0.060 vs. ≥ 0.805 is order-of-magnitude). Mendelian: HCCP is uniquely Pareto-frontier in a 2×2 window around the paper’s (0.85, 0.20) choice; $T_2 \geq 0.30$ lets B1 (gap 0.265) and wCP (gap 0.252) through, and $T_1 \geq 0.90$ or $T_2 \leq 0.15$ excludes HCCP itself — the operating window is narrow but the choice is not at the boundary.

Dataset	$T_1 \setminus T_2$	0.15	0.20	0.25	0.30
Complex ($K_{\text{eval}} = 5$)	≥ 0.80	H	H	H	H
	≥ 0.85	H	H	H	H
	≥ 0.90	H	H	H	H
Mendelian ($K_{\text{eval}} = 3$)	≥ 0.80	—	H	H	H, B1, wCP
	≥ 0.85	—	H	H	H, B1, wCP
	≥ 0.90	—	—	—	B1, wCP

851 The honest reading: on Complex, HCCP’s Pareto-uniqueness is robust to all reasonable threshold
852 choices we considered; the order-of-magnitude gap separation makes any threshold above $T_2 = 0.10$
853 select HCCP alone. On Mendelian, the threshold window is narrower — HCCP is unique on a 2×2
854 rectangle of $\{0.80, 0.85\} \times \{0.20, 0.25\}$, fails the floor at $T_2 \leq 0.15$ (HCCP’s own 0.173 gap), and
855 loses uniqueness at $T_2 \geq 0.30$ (where B1 split and wCP also satisfy the relaxed gap). The paper’s
856 $(0.85, 0.20)$ choice is at the lower-right corner of the unique region, not at its boundary; the claim
857 is not knife-edge but is dataset-asymmetric, consistent with the asymmetric H2H signal we report
858 throughout.

859 C.5 HCCP decoupled from the GBM aggregator (deferred from §6.2)

860 To rule out “HCCP’s coverage advantage rides on the GBM aggregator’s +14.8 pp AUPRC im-
861 provement”, we apply the same HCCP calibration on top of the published TraitGym LogReg ag-
862 gregator (Mendelian feature set CADD+GPN-MSA+Borzoi, AUPRC 0.752; Complex feature set
863 CADD+Borzoi, AUPRC 0.332) instead of our GBM. The variance head $\hat{\sigma}(x)$ is reused unchanged,
864 since it is fit on features X rather than on $\hat{p}_{\text{GBM}}$ (a per- x aleatoric estimator, base-predictor agnostic).
865 Tab. 14 reports $B = 200$ chrom-bootstrap CIs.

Table 14: HCCP applied on $(\hat{p}_{\text{LogReg}}, \hat{\sigma})$ vs. on $(\hat{p}_{\text{GBM}}, \hat{\sigma})$. The $\hat{\sigma}$ -bin gap and minority-class coverage are essentially unchanged across the base-predictor swap, confirming the conformal-side gain is decoupled from the AUPRC gain.

Dataset	Aggregator	Marg. cov.	Cov $ Y=1$	$\hat{\sigma}$ -bin gap $\downarrow$
Mendelian ($K_{\text{eval}} = 3$)	GBM (ours)	0.900 ± 0.012	0.875 ± 0.040	0.173 ± 0.192
	LogReg (TraitGym)	0.897 ± 0.018	0.897 ± 0.022	0.133 ± 0.197
Complex ($K_{\text{eval}} = 5$)	GBM (ours)	0.900 ± 0.003	0.905 ± 0.010	0.060 ± 0.031
	LogReg (TraitGym)	0.900 ± 0.004	0.901 ± 0.010	0.059 ± 0.034

866 The Complex row is essentially identical across the swap (gap 0.060 vs 0.059, $\text{cov}_{|Y=1}$ 0.905 vs
867 0.901). On Mendelian, HCCP-on-LogReg even gives a slightly tighter gap (0.133 vs 0.173) because
868 the LogReg point predictor is smoother and aligns more directly with the residual-trained $\hat{\sigma}$. Both
869 still clear the recommended ($\text{cov}_{|Y=1} \geq 0.85$, $\hat{\sigma}$ -bin gap ≤ 0.20) corner.

870 C.6 K_{eval} sensitivity sweep

871 Tab. 15 reports HCCP $\hat{\sigma}$ -bin gap against the strongest contemporary single-axis CP baseline (RLCP,
872 weighted CP, or SC-CP, whichever is tightest at each row) for $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$. HCCP partition
873 K is selected per outer fold by nested chrom-LOO; reporting metric is the worst-cell gap on a fixed
874 K_{eval} partition. Numbers are point estimates from a single seed (`T_tools/cp_baselines_h2h.py`
875 `-K-mode nested-cv`); the Mendelian $K_{\text{eval}} = 3$ row corresponds to Tab. 10’s 0.173 ± 0.192 HCCP
876 entry under bootstrap, and the small offset (0.163 here) reflects the bootstrap-vs-point-estimate
877 discrepancy.

878 C.7 Per-chromosome coverage

879 Tab. 16 reports per-chromosome coverage for the HCCP Mondrian predictor (CADD+GPN-
880 MSA+Borzoi, $\alpha = 0.10$, $K = 5$).

881 C.8 Per-trait clinical clustering (Complex traits)

882 To rule out “is HCCP holding because one easy trait class dominates the average”, we
883 cluster the 28 Complex traits and 30 Mendelian OMIMs into clinically-meaningful super-
884 categories and report per-cluster HCCP coverage. Cluster assignments are listed in
885 `T_tools/trait_cluster_aggregate.py`; weighted aggregation by trait sample size.

Table 15: K_{eval} sensitivity (single-seed point estimates). The recommended K_{eval} (Mendelian $K_{\text{eval}} = 3$, Complex $K_{\text{eval}} = 5$) keeps ≥ 100 minority-class samples per cell. **Complex:** HCCP dominates by $25\text{--}34\times$ at $K_{\text{eval}} \leq 5$ and by $3\text{--}4\times$ at $K_{\text{eval}} \geq 8$, robustly across the sweep. **Mendelian:** HCCP wins by $1.15\times$ at the recommended $K_{\text{eval}} = 3$ and ties at $K_{\text{eval}} = 2$; at $K_{\text{eval}} \geq 5$, the per-cell minority count drops below ~ 70 and weighted CP / RLCP catch up because the heteroscedastic head can no longer be reliably calibrated on the smaller per-cell partition. The Mendelian-specific narrowing of margin at high K_{eval} is consistent with the per-cell sample-size floor argument that motivates the recommended K_{eval} in the first place (§5.2).

Dataset	K_{eval}	HCCP gap	RLCP	wCP	SC-CP	ratio (best baseline / HCCP)
Mendelian	2	0.233	0.233	0.233	0.233	$1.00\times$
	3 (recommended)	0.163	0.321	0.188	0.374	$1.15\times$ (vs wCP)
	5	0.567	0.483	0.379	0.567	$0.67\times$ (HCCP loses to wCP)
	8	0.700	0.600	0.613	0.700	$0.86\times$ (HCCP loses to RLCP)
	10	0.700	0.600	0.613	0.700	$0.86\times$ (HCCP loses to RLCP)
Complex	2	0.027	0.799	0.803	0.803	$30.05\times$
	3	0.032	0.807	0.807	0.807	$25.39\times$
	5 (recommended)	0.024	0.802	0.802	0.802	$33.80\times$
	8	0.268	0.831	0.831	0.831	$3.10\times$
	10	0.208	0.847	0.847	0.847	$4.08\times$

Table 16: Per-chromosome coverage under Mondrian- $(y \times \hat{\sigma}\text{-bin})$, Mendelian dataset. Target: 0.90. Chromosomes 17–22, X are test chroms under chrom-LOO. Worst per-chrom $|\text{gap}| = 0.100$ at chr6 (cov 0.800); the second-worst is chr19 ($|\text{gap}| = 0.073$). The cell-level KS audit (Tab. 4) shows that this cross-chromosome variability is dataset-intrinsic on Mendelian; T3’ supplies the certificate (App. A.7).

Chrom	n	Coverage	$ \text{gap} $
1	210	0.857	0.043
2	230	0.896	0.004
3	155	0.945	0.045
5	100	0.850	0.050
6	150	0.800	0.100
7	210	0.943	0.043
8	140	0.943	0.043
9	240	0.842	0.058
10	190	0.879	0.021
11	160	0.944	0.044
12	150	0.933	0.033
13	210	0.848	0.052
14	100	0.900	0.000
16	80	0.913	0.013
17	150	0.900	0.000
19	200	0.828	0.073
20	50	0.900	0.000
22	20	0.950	0.050
X	250	0.958	0.058
Overall	3380	0.902	0.002

886 C.9 Cross-dataset transfer

887 **Directional asymmetry.** C→M transfer preserves coverage (0.901) because the complex-trait
888 feature space subsumes the monogenic signal. M→C collapses (0.738) because monogenic variants
889 occupy a narrow subspace and the base predictor itself loses discriminative signal (§7, R3). The
890 Barber T4 proxy bound (0.742) is tight for M→C (observed 0.738, within 0.004) and conservative
891 for C→M.

HCCP per-chromosome coverage --- $K = 5$, $\alpha = 0.10$, target 0.90. Marker size $\propto \log n_c$; red dashed = T3' robust lower bound.




Figure 2: Per-chromosome HCCP coverage on chrom-LOO, hybrid forest layout ($\alpha = 0.10$, $K = 5$, target 0.90). Each row is one chromosome (sorted ascending by HCCP coverage); the horizontal bar is the 95% within-chrom bootstrap CI ($B = 1000$); the marker is the HCCP point estimate, with **marker size** $\propto \log n_c$ so that small chromosomes carry small markers and (by the $\text{Bin}(n_c, 0.90)$ floor) wider CIs. Light green shading: ± 0.03 band around target; black dashed line: $1 - \alpha = 0.90$. Red dashed line: T3' robust lower bound ≥ 0.59 , derived from the Mendelian worst-cell $\bar{\delta}_{\text{TV}}^{(k,b)} \leq 0.31$ via Theorem 2. On **Mendelian**, the worst chromosome chr6 (cov = 0.80, $n = 30$) has the smallest marker and the widest CI $[0.63, 0.93]$ that *crosses* the target — its under-coverage is not statistically distinguishable from binomial noise. All chromosomes on both datasets remain above the T3' certificate, and on **Complex** all 22 chromosomes' CIs lie within ± 0.04 of target. The marker-size encoding makes the T3' regime visually self-evident: small markers concentrate at the low-coverage tail, confirming that cross-chromosome variability is a sample-size effect (the variance term of Eq. 3) rather than a method failure. The companion table 16 reports the per-chrom numerics.

Table 17: Per-cluster HCCP coverage on the 28 Complex traits with ≥ 10 positives (trait-LOO, CADD+GPN-MSA+Borzoï, $\alpha = 0.10$, $K = 2$). All three clusters hold target marginal coverage to within ± 0.005 and minority-class coverage to within ± 0.04 ; the per-trait worst gap stays ≤ 0.043 in each cluster, confirming that the trait-LOO floor in the headline is not driven by one easy cluster.

Cluster	# traits	n_{tot}	Marg. cov.	Cov Y=1	Worst per-trait gap $\downarrow$
Cardio-metabolic / anthropometric	17	3940	0.905	0.906	0.033
Hematologic / blood-cell	8	2160	0.896	0.912	0.043
Inflammatory / pulmonary / other	3	600	0.905	0.933	0.017
Overall (28 traits)	28	6700	0.902	—	0.043

892 C.10 K-sweep ablation

893 Tab. 20 validates Theorem 3 (T5.1) empirically by sweeping $K \in \{2, 3, 5, 8, 10, 15, 20, 30\}$. For each
894 dataset the plug-in Lipschitz estimate is $\hat{L}_F \approx 49.4$ (Mendelian) and 21.8 (Complex) via 100-batch
895 finite differences on $\hat{\sigma}(X)$; the resulting oracle bin count $K^* = \lfloor \sqrt{L_F R \pi_{\min} n} \rfloor$ is 88 (Mendelian)
896 and 110 (Complex), neither of which is realizable because the finite-sample variance term $K/(\pi_{\min} n)$
897 dominates long before the stationarity bias term $L_F R/K$ crosses it. The bolded rows in Tab. 20
898 mark the K that minimises the chrom-LOO test-fold worst-cell gap ($K = 3$ Mendelian, $K = 2$
899 Complex) under the *legacy non-nested* K-selection protocol; the proper nested chrom-LOO selector
900 (§5.2) instead picks the modal $\hat{K}(c_{\text{outer}}) \in \{5, 8\}$ on Mendelian and $\hat{K}(c_{\text{outer}}) = 5$ on Complex,
901 restoring CP exchangeability. The fallback column reports the number of (k, b) cells that triggered
902 pooled class-conditional backoff ($n_{k,b} < 5$, §4).

903 **Reading the table — three regimes.** The K-sweep extended to $K = 200$ exposes *three* regimes
904 of HCCP behaviour, only the first two of which are predicted by Theorem 3. (i) **Bias-dominated**

Table 18: Per-cluster HCCP coverage on the 30 Mendelian OMIMs with ≥ 5 positives (trait-LOO, CADD+GPN-MSA+Borzoi, $\alpha = 0.10$, $K = 3$). Skeletal/connective shows the highest worst per-trait gap (0.122) driven by under-coverage of the minority pathogenic class ($\text{cov}_{|Y=1} = 0.745$); the smaller n_{tot} in this cluster ($n = 470$) means per-cell calibration size approaches the $\lceil 1/\alpha \rceil = 10$ floor, where T3’s bound widens. The other two clusters hold target tightly.

Cluster	# OMIMs	n_{tot}	Marg. cov.	$\text{Cov}_{ Y=1}$	Worst per-trait gap $\downarrow$
Skeletal / connective	8	470	0.913	0.745	0.122
Neuromuscular / nervous-system	10	610	0.897	1.000	0.067
Metabolic / hematologic / other	12	1340	0.909	0.918	0.100
Overall (30 OMIMs)	30	2420	0.907	—	0.122

Table 19: Cross-dataset transfer (CADD+Borzoi). M→C: train on Mendelian, test on Complex. C→M: reverse. AUPRC chrom-weighted (TraitGym convention). Marginal-gap = marginal coverage minus target 0.90 (signed). T4 proxy bound = $1 - \alpha - d_{\text{TV}}$.

Direction	AUPRC _{chrom}	Marg. cov.	Marg. gap	TV proxy	T4 bound
M→C	0.205	0.738	−0.162	0.158	0.742
C→M	0.665	0.901	+0.001	0.158	0.742

905 ($K \leq 3$ Mendelian, $K \leq 5$ Complex): worst-cell gap falls as K increases because within-bin
906 score heterogeneity decreases. Sweet spot under legacy chrom-LOO test-gap minimisation is $K=3$
907 (Mendelian) and $K=2$ (Complex); the rectified nested-CV $\hat{K}(c_{\text{outer}})$ is larger ($\{5, 8\}$ Mendelian, 5
908 Complex) because the inner-fold scoring is on a fixed K_{eval} axis (§5.2). (ii) **Variance-dominated**
909 ($K \in [5, 30]$ Mendelian, $K \in [10, 30]$ Complex): worst-cell gap rises as cells drop toward the
910 $n_{k,b} \geq \lceil 1/\alpha \rceil = 10$ floor. Mendelian saturates at gap 0.73; Complex at 0.60. This is the tradeoff
911 captured by Eq. 3. (iii) **Fallback-dominated** ($K \geq 50$, not modeled by T5): once the majority of
912 cells trigger the pooled-fallback ($n_{k,b} < 5$), the predictor effectively reverts to T2 class-conditional
913 calibration, which is uniform across $\hat{\sigma}$ -bins. Worst-cell gap therefore stops growing and partially
914 recovers (Mendelian: $0.73 \rightarrow 0.10$ as K grows from 30 to 200; Complex: 0.90 plateau and 0.70).
915 This third regime is invisible in the asymptotic T5 bound but visible at the realistic $n \leq 11,400$
916 scale of TraitGym, and it explains why both the legacy non-nested $K = 2, 3$ and the nested-CV
917 modal $\hat{K} = 5-8$ remain far below the oracle K^* derived from any L_F estimator in Tab. 21: the
918 variance-dominated regime never “wins” before fallback takes over.

919 **L_F estimator with concentration bound.** The oracle bin count $K^* = \lfloor \sqrt{L_F R \pi_{\min} n} \rfloor$ depends
920 on the plug-in Lipschitz estimate $\hat{L}_F$. We adapt the Lipschitz-Constant Least-Squares estimator
921 of Huang et al. [2024] to our score-CDF Lipschitz problem: bin $\hat{\sigma}$ into 30 quantile bins per class
922 with at least 20 variants per bin, compute pairwise Kolmogorov–Smirnov distances and $\hat{\sigma}$ -gaps, then
923 regress $\text{KS}_{ab} = \hat{L}_F \cdot \text{gap}_{ab}$ via least squares through the origin. Bootstrap with $B = 500$ chrom-level
924 resamples gives a 95% CI. Tab. 21 shows that the LCLS estimator yields tight CIs ($\hat{L}_F = 2.97$
925 Mendelian, 4.51 Complex), an order of magnitude smaller than the legacy max-of-ratios estimator
926 (183 and 53.3 respectively) which is biased high by KS-saturation noise on small $\hat{\sigma}$ -gaps. Substituting
927 LCLS $\hat{L}_F$ into Theorem 3 gives $K^* = 22$ (Mendelian) and $K^* = 50$ (Complex), still above the
928 modal nested-CV $\hat{K}(c_{\text{outer}}) \in \{5, 8\}$ on Mendelian and $\hat{K}(c_{\text{outer}}) = 5$ on Complex, by a factor of
929 $3-10\times$. The remaining gap is the finite-sample fallback effect (App. C.10, third regime).

930 C.11 DEGU-lite vs HCCP head-to-head

931 Tab. 22 compares HCCP’s supervised $\hat{\sigma}$ against DEGU-lite’s ensemble $\hat{\sigma}$ under identical Mondrian-
932 ($y \times \hat{\sigma}$ -bin) conformal calibration and feature sets.

933 **Interpretation.** On Mendelian (monogenic, strong signal), ensemble disagreement is a reasonable
934 uncertainty proxy: DEGU-lite achieves slightly lower worst-cell gap (0.033 vs. 0.041). On Complex
935 (polygenic, dominant aleatoric noise), supervised residual regression provides $4.3\times$ better worst-cell
936 coverage (0.013 vs. 0.057). This confirms the epistemic–aleatoric decomposition discussed in §7:

Table 20: K -sweep on Mendelian and Complex (CADD+GPN-MSA+Borzoi, $\alpha = 0.10$, single seed, legacy non-nested K -selection). Worst-cell gap = $\max_{k,b} |\text{cov}_{k,b} - 0.90|$ on the test fold; $n_{\min}$ = minimum calibration count per cell; fallback = cells triggering pooled backoff. Bolded rows are the legacy minimum (test-fold leakage; superseded by nested-CV $\hat{K}(c_{\text{outer}})$ in §5.2 and Tab. 21). Single-seed values here are smaller than the bootstrap means in Tab. 10 because they ride the favourable test fold the legacy selector tunes against.

K	Mendelian				Complex			
	Worst gap	Mean gap	$n_{\min}$	fallback	Worst gap	Mean gap	$n_{\min}$	fallback
2	0.233	0.059	6	0	0.006	0.002	237	0
3	0.047	0.012	19	1	0.010	0.004	131	0
5	0.285	0.041	13	2	0.042	0.007	52	0
8	0.627	0.067	11	5	0.058	0.010	19	0
10	0.700	0.059	10	6	0.208	0.017	13	0
15	0.733	0.049	6	9	0.600	0.034	10	0
20	0.733	0.047	6	14	0.600	0.030	10	1
30	0.344	0.031	5	22	0.567	0.034	9	2
<i>Fallback-dominated regime: $n_{k,b} < \lceil 1/\alpha \rceil = 10$ triggers pooled class-conditional threshold</i>								
50	0.264	0.038	5	39	0.900	0.094	5	13
100	0.186	0.052	5	96	0.900	0.080	5	92
200	0.100	0.046	5	188	0.700	0.051	5	245

Table 21: Four L_F estimators and their oracle $K^* = \lfloor \sqrt{L_F R \pi_{\min} n} \rfloor$ and $G(K^*) \leq 2\sqrt{L_F R / (\pi_{\min} n)}$ bounds. CIs are 95% bootstrap percentile intervals over $B = 500$ chrom-level resamples. The modal nested-CV $\hat{K}(c_{\text{outer}})$ (§5.2) is reported for comparison; the legacy non-nested K_{CV} (test-gap argmin) values 3 Mendelian and 2 Complex are deprecated. [†]On Complex, the full-data legacy point estimate 53.3 lies just below CI lower bound: under bootstrap resampling the legacy max-of-ratios is consistently larger (median 72.1, mean 73.8), reflecting its instability under repeated extreme-pair concentration.

Estimator	Dataset	$\hat{L}_F$	95% CI	K^*	$G(K^*)$ bound
LCLS regression [Huang et al., 2024]	Mendelian	2.97	[2.89, 3.45]	22	0.130
isotonic envelope	Mendelian	26.9	[6.85, 93.2]	61	0.388
legacy max-of-ratios	Mendelian	183	[138, 299]	169	1.011
<i>modal nested-CV $\hat{K}(c_{\text{outer}})$ (Mendelian)</i>				5–8	0.100 (observed)
LCLS regression [Huang et al., 2024]	Complex	4.51	[4.47, 4.69]	50	0.088
isotonic envelope	Complex	8.02	[4.48, 33.1]	67	0.118
legacy max-of-ratios	Complex	53.3	[54, 103] [†]	173	0.303
<i>modal nested-CV $\hat{K}(c_{\text{outer}})$ (Complex)</i>				5	0.024 (observed)

937 where model disagreement and data noise contribute equally, both $\hat{\sigma}$ sources suffice; where data noise
938 dominates (Complex traits), the supervised head captures what ensembles cannot.

939 C.12 A-SL audit (deferred from §3.1)

940 C.13 $\hat{\sigma}$ class ablation: GBM vs. MLP vs. Random Forest

941 T5.1’s dimension-free $O(n^{-1/2})$ rate (Theorem 3) builds on 4: $\hat{\sigma}$ is a 1-D Lipschitz summary with
942 $L_F < \infty$. The headline LCLS estimate $\hat{L}_F^{\text{GBM}} = 2.97/4.51$ (Mendelian / Complex) was computed
943 on the single GBM-based $\hat{\sigma}$ used throughout the paper. Reviewers reasonably ask whether the
944 dimension-free property depends on a smoothness property of the GBM architecture rather than a
945 property of the conformal construction.




Figure 3: Three L_F estimators paired with their implied asymptotic K^* , with the empirical modal nested-CV $\hat{K}(c_{outer})$ overlaid as the operative reference. **Left:** forest of $\hat{L}_F$ on log scale, 95% bootstrap CIs ($B = 500$ chrom-resamples); marker size encodes precision ($1/\log_{10}$ of CI ratio). LCLS yields a tight CI ($1.19\times$ on Mendelian, $1.05\times$ on Complex); isotonic envelope is order-of-magnitude wider; legacy max-of-ratios sits at the KS-statistic noise floor. **Right:** horizontal bar of the implied oracle $K^* = \lfloor \sqrt{L_F R \pi_{\min} n} \rfloor$ (linear axis); the red dashed band/line marks the modal nested-CV $\hat{K}(c_{outer}) \in \{5, 8\}$ (Mendelian) and $= 5$ (Complex) per §5.2. **Punchline:** every asymptotic K^* exceeds $\hat{K}_{\text{modal}}$ by 3–10 $\times$ on the trustworthy LCLS estimator and by orders of magnitude on the loose ones, identifying the App C.10 fallback-dominated regime as the binding finite-sample constraint at $n \leq 11,400$. On Complex, the legacy marker uses bootstrap median 72.1 rather than the unstable full-data point 53.3 (see Tab. 21 footnote).

Table 22: HCCP vs DEGU-lite: same features (CADD+GPN-MSA+Borzoi), same Mondrian partition at $K = 5$, different $\hat{\sigma}$ source. Worst gap = max per- $\hat{\sigma}$ -bin marginal coverage deviation from target 0.90 (the same metric as Tab. 2, smaller than the cell-level worst- (k, b) gap of Tab. 9). Single-seed point estimates.

$\hat{\sigma}$ source	Dataset	Cov	Cov $ _{Y=1}$	Worst gap	Mean gap
HCCP (learned)	Mendelian	0.902	0.905	0.041	0.022
DEGU (ensemble)	Mendelian	0.901	0.917	0.033	0.020
HCCP (learned)	Complex	0.901	0.908	0.013	0.007
DEGU (ensemble)	Complex	0.901	0.907	0.057	0.028

To stress-test this, we re-train two alternative $\hat{\sigma}$ heads on identical features (CADD+GPN-MSA+Borzoi, $d \approx 8500$), identical chrom-LOO protocol, and identical Day 10 GBM aggregator $\hat{p}$:

- **MLP:** 3-layer 64-32-1 with softplus output and AdamW training on the absolute-residual target (80 epochs, lr $5 \cdot 10^{-3}$, weight decay 10^{-4} , batch 256, dropout 0.1).
- **Random Forest:** 50-tree ensemble with max_features="sqrt" (a non-smooth, piecewise-constant $\hat{\sigma}$ used as the worst-case stress test for 4).

Both heads share the same supervised target $|r_i| = |y_i - \hat{p}^{(-C_i)}(x_i)|$ as the GBM head. Per-head LCLS $\hat{L}_F$ (regression slope of pairwise KS on σ -gap), HCCP single-seed worst-cell $\hat{\sigma}$ -bin gap at the recommended K_{eval} , and Spearman($\hat{\sigma}, |r|$) are reported below.

Reading the ablation. On Complex ($n = 11,400$, per-cell minority ≈ 200): all three $\hat{\sigma}$ classes give worst-cell gap $\in [0.011, 0.015]$ and minority-coverage within ± 0.005 of target. $\hat{L}_F$ varies by less than $2\times$ across classes (2.59 to 4.67). The dimension-free claim of T5.1 is empirically robust to $\hat{\sigma}$ architecture once n is large enough to keep per-cell minority above the variance-stable floor.

Bootstrap robustness ($B = 200$ chromosome-level resamples) of HCCP on the $(\text{Cov}|Y=1, \hat{\sigma}\text{-bin gap})$ plane




Figure 4: Joint bootstrap robustness of HCCP on the $(\text{Cov}|Y=1, \hat{\sigma}\text{-bin gap})$ plane (the same axes as Fig. 1(b)): $B = 200$ chromosome-level resamples at $K = 5$, $\alpha = 0.10$. Each panel shows replicate scatter (red points), 50% / 95% Gaussian-KDE contours, the bootstrap mean (+), and the single-seed all-chrom point estimate (★, gold). Marginal KDEs sit on the top and right axes; dashed lines mark the all-chrom point on each marginal. **Mendelian** (left): bootstrap 95% CIs $\text{Cov}|Y=1 \in [0.858, 0.918]$, $\hat{\sigma}\text{-bin gap} \in [0.004, 0.044]$; the all-chrom point (0.905, 0.005) lies inside the 50% contour. The wider cloud reflects chrom-level heterogeneity (Fig. 2). **Complex** (right): $\text{Cov}|Y=1 \in [0.894, 0.911]$, $\text{gap} \in [0.002, 0.006]$; the all-chrom point (0.908, 0.002) sits at the cloud centre. On both datasets the bootstrap CIs lie strictly inside the $(\text{Cov}|Y=1 \geq 0.85, \text{gap} \leq 0.20)$ ideal corner of Fig. 1(b), demonstrating that the headline frontier-point claim is not driven by a single favourable chromosome split.

Table 23: $\hat{\sigma}$ class ablation. Single-seed point estimates from `T_tools/sigma_ablation.py`; cross-method-fair worst-cell gap with min-cell-count 5 (paper convention, Tab. 10). L_F^{LCLS} is the regression slope of pairwise KS on $\hat{\sigma}$ -gap; larger L_F indicates a less smooth $\hat{\sigma} \rightarrow F_{k,\sigma}$ map. Spearman quantifies how well $\hat{\sigma}$ tracks the magnitude of the actual residual $|y - \hat{p}(x)|$ on chrom-LOO.

Dataset	$\hat{\sigma}$ class	$\hat{L}_F^{\text{LCLS}}$	Worst-cell gap	$\text{Cov} Y=1$	Spearman($\hat{\sigma}, r $)
Mendelian ($n = 3,380$)	GBM (paper)	2.97	0.047	0.891	0.754
	MLP	73.96	0.017	0.888	0.125
	Random Forest	2.31	0.900	0.876	0.575
Complex ($n = 11,400$)	GBM (paper)	4.51	0.015	0.905	0.689
	MLP	2.59	0.011	0.904	0.630
	Random Forest	4.67	0.011	0.903	0.605

960 **On Mendelian** ($n = 3,380$, per-cell minority $\in [3, 316]$): the worst-cell gaps span three orders of
961 magnitude (0.017, 0.047, 0.900). Two diagnostics resolve this:

- 962 • The MLP’s nominal $\hat{L}_F = 74$ is an artifact of $\hat{\sigma}$ -collapse: the 64-channel projection cannot
963 capture the $d \approx 8500$ feature space at $n = 3,380$, and the MLP’s output collapses to
964 near-zero ($\hat{\sigma}$ q05/q95 = $4 \cdot 10^{-5} / 3 \cdot 10^{-5}$, Spearman = 0.125). The pairwise σ -gaps in the
965 LCLS estimator collapse to numerical floor, inflating the KS/gap ratio. Operationally HCCP
966 degenerates to class-conditional split CP via the small-cell pooled fallback (Alg. 1, line 6);
967 the gap of 0.017 is a class-cond split CP property, not an HCCP property.
- 968 • The Random Forest’s nominal worst-cell gap of 0.900 is driven by a single minority cell
969 ($y = 1, b = 0, n = 5$, coverage = 0/5): RF’s heavy-tailed $\hat{\sigma}$ distribution puts 5 minority
970 variants in the lowest $\hat{\sigma}$ -bin where the per-chrom calibration quantile is misspecified. The
971 second-worst cell of Mendelian/RF has gap 0.013. With a stricter min-cell threshold (≥ 10)
972 the worst-cell gap drops to 0.013, comparable to GBM. Sample-size noise drives the 0.900
973 figure, not a violation of 4.




Figure 5: HCCP vs DEGU-lite worst $\hat{\sigma}$ -bin coverage gap, paired by dataset (Mendelian top, Complex bottom). Each dumbbell connects HCCP (blue circle, learned residual $\hat{\sigma}$) and DEGU-lite (red square, ensemble-distilled $\hat{\sigma}$) at matched GBM backbone and matched HCCP-style Mondrian calibration; mid-line annotation reports the head-to-head ratio and winner. The green band marks the operational ideal zone ($\text{gap} \leq 0.03$). On **Complex** (aleatoric-dominated) the dumbbell stretches across the full panel: HCCP 0.013 vs. DEGU 0.057 ($4.33\times$, HCCP wins, only HCCP inside the ideal zone). On **Mendelian** (monogenic, strong signal) the dumbbell is short: DEGU 0.033 vs. HCCP 0.041 ($1.23\times$, DEGU marginally lower); both points sit just outside the ideal zone, an honest near-tie. Marginal coverage and per-class coverage are reported in Tab. 22 above; singleton-fraction efficiency at the recommended operating point appears in Tab. 2 (HCCP 0.776 Mendelian / 0.228 Complex). The figure focuses on the headline coverage metric where the $\hat{\sigma}$ -source choice binds. The asymmetry — decisive HCCP win on aleatoric-dominated Complex, near-tie on Mendelian where ensemble disagreement is an adequate proxy — is the App C reading of the epistemic–aleatoric decomposition in §7.

974 The Mendelian discrepancies are consistent with Eq. (3): the variance term $K/(\pi_{\min}n)$ dominates at
 975 $n = 3,380$, $K = 3$, $\pi_{\min} = 0.10$, putting the empirical worst-cell gap floor at $\sim K/(\pi_{\min}n) \approx 0.009$
 976 even before any $\hat{\sigma}$ -related bias, plus a small-cell binomial tail that adds variant-count outliers. T5.1’s
 977 upper bound $G(K^*) \leq 2\sqrt{L_F R/(\pi_{\min}n)}$ remains valid for all three $\hat{\sigma}$ classes, but at small n the
 978 empirical gap is dominated by the small-cell variance regime that $\hat{\sigma}$ choice cannot cure — precisely
 979 the regime boundary already documented in §6.2 (R1) and Tab. 15. **We conclude: the dimension-**
 980 **free property of T5.1 is architecture-agnostic on Complex (where the rate regime binds); on**
 981 **Mendelian, the small-cell variance regime takes over, and GBM remains the operational choice**
 982 **given its Spearman fit (less small-cell noise) at finite L_F .**

983 D α sensitivity

984 Figure 7 and Tab. 24 report HCCP’s empirical coverage and $\hat{\sigma}$ -bin gap as the miscoverage level α
 985 sweeps over $\{0.01, 0.05, 0.10, 0.15, 0.20, 0.30, 0.50\}$, fixed feature set CADD+GPN-MSA+Borzoï,
 986 fixed $K = 5$, single seed. The target $1 - \alpha$ line is included for reference.

987 **Observations.** (i) The marginal coverage is calibrated to within ≤ 0.015 of target at every α ,
 988 consistent with the exact T1 two-sided bound $[1 - \alpha, 1 - \alpha + 1/(n_{\min} + 1)]$. (ii) The worst-cell
 989 $\hat{\sigma}$ -bin gap stays below the $\alpha = 0.10$ headline 0.04 threshold across the entire α range, confirming the
 990 T3 bound is not an artifact of a specific α choice. (iii) The singleton fraction grows monotonically
 991 with α on Complex (from 3% at $\alpha = 0.01$ to 77% at $\alpha = 0.50$) but is non-monotone on Mendelian
 992 (peaks at 82% near $\alpha = 0.15$, falls back at higher α as empty sets $\emptyset$ emerge); this is expected for
 993 class-imbalanced problems where a very loose α can drop a small- $\hat{\sigma}$ variant below the lowered
 994 quantile bar on both classes simultaneously.

995 **Operating-point recommendations.** The singleton fraction column of Tab. 24 is the practical
 996 knob: low α buys high coverage at the cost of frequent abstention; high α produces more confident
 997 calls at the cost of more uncovered cases. We provide three benchmarked operating points:




Figure 6: **A-SL audit**. Per-pair Lipschitz ratio $\rho_{k,a,b} = \text{KS}(\hat{F}_{k,a}, \hat{F}_{k,b}) / |\bar{\sigma}_a - \bar{\sigma}_b|$ on $\hat{\sigma}$ -bin pairs. Noise-adjusted upper envelope is finite without fat-tail breakdown: $L_F^{(95\%)} = 13.8$ (Mendelian) / 8.0 (Complex). LCLS estimator $\hat{L}_F \approx 3/4.5$ measures the typical local slope and underestimates the upper envelope by $\sim 3\text{--}5\times$.

Table 24: α sweep: marginal coverage, $\hat{\sigma}$ -bin gap, and singleton fraction (Mendelian / Complex, CADD+GPN-MSA+Borzoï, $K = 5$, seed 42). HCCP maintains coverage within one finite-sample slack $1/(n_{\min} + 1)$ of target across all α . Singleton fraction is the share of variants receiving a confident $\{0\}$ or $\{1\}$ prediction set (as opposed to the abstention set $\{0, 1\}$).

α	M cov	M $\hat{\sigma}$ -gap	C cov	C $\hat{\sigma}$ -gap	singleton (M / C)
0.01	0.991	0.010	0.991	0.002	53% / 3%
0.05	0.950	0.019	0.951	0.001	68% / 11%
0.10	0.902	0.009	0.901	0.001	80% / 22%
0.15	0.851	0.006	0.852	0.002	82% / 34%
0.20	0.803	0.012	0.802	0.002	79% / 45%
0.30	0.704	0.013	0.701	0.003	70% / 67%
0.50	0.507	0.009	0.501	0.003	52% / 77%

- 998 • **High-stakes rare-disease classification** ($\alpha = 0.10$): Mendelian singleton 80%, Complex
999 singleton 22%. Use when a missed pathogenic variant has clinical-grade cost (e.g., hereditary
1000 cancer panels). The Complex 22% rate signals that polygenic variants are not individually
1001 classifiable at this confidence; for population genetics this is the wrong operating point.
- 1002 • **Polygenic screening triage** ($\alpha = 0.20$): Mendelian singleton 79%, Complex singleton
1003 45%. Matches the operating point of Adegoké et al. [2025] for clinical conformal-prediction
1004 deployment (41% abstention reported). Use for prioritising variants for follow-up genotyping
1005 or family studies.
- 1006 • **Hypothesis generation** ($\alpha = 0.30$): Mendelian singleton 70%, Complex singleton 67%.
1007 Use when downstream review is cheap and the goal is to widen the candidate set (e.g.,
1008 GWAS fine-mapping, drug-target prioritisation).




Figure 7: Compact 1×3 α -sweep with datasets encoded by line style: *Mendelian* = solid line, filled markers; *Complex* = dashed line, hollow markers. (a) Calibration residual $\text{cov} - (1 - \alpha)$ for Split CP (homoscedastic class-cond, gray), HCCP class-cond (blue), and HCCP Mondrian ($y \times \hat{\sigma}$ -bin, $K=5$) (red); shaded band is the finite-sample slack $\pm 1/(n_{\min}+1)$. All six (method $\times$ dataset) curves track target to within the slack at every α . (b) $\hat{\sigma}$ -bin coverage range $\max_b \text{cov}_b - \min_b \text{cov}_b$ for HCCP Mondrian, grouped per α (Mendelian red bars left, Complex blue bars right; numeric annotations on small Complex bars); dashed reference is the $\alpha=0.10$ headline threshold ≤ 0.04 , respected at every α on both datasets. (c) Singleton fraction (the operational knob); the three benchmarked operating points $\alpha \in \{0.10, 0.20, 0.30\}$ are marked by vertical guides and labeled at top — *high-stakes / triage / hypothesis* (full operating-point recommendations in the body text below). The arrow callout marks the Mendelian non-monotonicity peak (≈ 0.82 at $\alpha = 0.15$): at large α on the imbalanced minority-class regime, empty-set abstentions $\emptyset$ replace singletons, an honest finite-sample artefact.

1009 The coverage guarantee remains $1 - \alpha$ at each operating point; only the singleton/abstention budget
1010 shifts. This is the key practical advantage of conformal prediction over a fixed-threshold classifier:
1011 the user, not the model designer, owns the coverage–decisiveness tradeoff.

1012 E Cross-domain stress test — ProteinGym

1013 To test whether HCCP’s coverage guarantees generalize *beyond* human variant pathogenicity clas-
1014 sification, we evaluate on **ProteinGym** [Notin et al., 2023], a large-scale protein variant fitness
1015 benchmark with 217 deep mutational scanning (DMS) assays across 2.4M mutants. We use the
1016 DMS_score_bin discretization provided by the benchmark authors ($\approx 57\%$ positive, high-fitness
1017 mutations) as the binary target, making the task directly comparable to HCCP’s TraitGym setting.

1018 **Features.** Unlike TraitGym, ProteinGym has no pre-computed aggregator-ready features — the
1019 benchmark’s baselines use pretrained language models (ESM-1v, MSA Transformer, Tranception).
1020 To keep this cross-domain test CPU-only and independent of any particular foundation model, we
1021 engineer a minimal 67-dimensional feature set from the mutation specification alone (App. E.1):
1022 one-hot encodings of wild-type and mutant amino acids, position normalized by sequence length,
1023 BLOSUM62 substitution score, physicochemical deltas (hydrophobicity, volume, charge, polarity),
1024 and local composition in a ± 5 AA window. These features are protein-agnostic, so a single classifier
1025 trained on non-target assays should generalize structurally across new proteins; this is precisely the
1026 assumption HCCP’s cross-assay calibration tests.

1027 **Protocol (assay-LOO).** For each target assay, we train the base classifier $\hat{p}$ on all mutants from the
1028 remaining 216 assays, fit $\hat{\sigma}$ on OOF residuals using the same pool, and calibrate Mondrian- $(y \times \hat{\sigma}$ -bin)
1029 conformal thresholds ($K = 5$, $\alpha = 0.10$, minimum cell size 10). We then compute marginal coverage,
1030 $\hat{\sigma}$ -bin coverage gap, and singleton fraction on the target assay. This is the ProteinGym analogue of
1031 the trait-LOO stress test reported in §6.

1032 **Results.** Across 50 assay-LOO held-out proteins (each trained on the union of the remaining
1033 $\approx 357k$ mutants from 166 other assays), HCCP’s marginal coverage tracks the target to within one
1034 finite-sample slack, and the $\hat{\sigma}$ -bin gap distribution is tight and asymmetric around the TraitGym
1035 chrom-LOO reference.

Table 25: ProteinGym assay-LOO: HCCP coverage statistics aggregated over the 50 largest valid assays (≥ 200 mutants, ≥ 50 positives; per-assay mutants capped at 3,000). $\alpha = 0.10$, $K = 5$. Source: `outputs/proteingym_hccp_n50/summary.json`.

Metric	Value
AUPRC test (mean over assays)	0.709
AUPRC test (median)	0.678
Marginal coverage (mean)	0.887
Marginal coverage (std)	0.028
Coverage within ± 0.02 of target	27/50 assays
Coverage within ± 0.05 of target	44/50 assays
$\hat{\sigma}$ -bin gap (mean)	0.052
$\hat{\sigma}$ -bin gap (median)	0.042
$\hat{\sigma}$ -bin gap (p90)	0.095
Frac. singleton (mean)	0.358




Figure 8: ProteinGym cross-domain stress test (50 held-out assays). **Left:** marginal coverage vs. test-assay AUPRC per held-out protein; the 6 under-covering assays ($\text{cov} < 0.85$) are coloured red against the 44 in-band assays in grey. Coverage stays within ± 0.02 of the 0.90 target for 27/50 assays and within ± 0.05 for 44/50. AUPRC and coverage are positively correlated (Spearman $\rho = 0.46$, $p = 0.001$): when the base predictor is strong, coverage tracks the target. **Right:** distribution of $\hat{\sigma}$ -bin coverage gap across the 50 assays (marginal-bin range, same metric as Tab. 2), with the ProteinGym assay mean (≈ 0.052) and the TraitGym Mendelian marginal-bin reference (0.020, Tab. 2) overlaid. The ProteinGym distribution sits at $\sim 2\times$ the TraitGym Mendelian marginal-bin reference (median 0.042, p90 0.095), well below the cell-level worst- (k, b) ceiling of 0.173 (Tab. 10), consistent with a bounded cross-domain shift.

1036 **Interpretation against pre-specified predictions.** We pre-specified three reading criteria in the
1037 proposal submission; they evaluate as follows. (i) *Coverage within ± 0.02 of target.* **Achieved on the**
1038 **aggregate:** mean coverage 0.887 across 50 held-out assays, $|\text{gap}| = 0.013$, std across assays 0.028.
1039 Per-assay, 27/50 assays hit the ± 0.02 band and 44/50 hit the ± 0.05 band, with 6/50 outliers below
1040 0.85 documented in the failure-mode analysis below. HCCP’s marginal coverage holds within ± 0.05
1041 of target on the majority of held-out assays whose sequence is unseen at training, but degrades on
1042 assays where the base predictor is weak (Spearman $\rho_{\text{AUPRC}, \text{cov}} = 0.46$). (ii) *$\hat{\sigma}$ -bin gap at TraitGym*
1043 *trait-LOO scale* (0.002–0.004). **Not achieved:** ProteinGym gap has median 0.042, roughly an order
1044 of magnitude above the TraitGym trait-LOO floor. We interpret this as an *honest OOD signal*: each
1045 new protein has assay-specific catalytic residues, structural hotspots, and readout idiosyncrasies (e.g.,
1046 surface expression vs. activity in the CP2C9 pair) that physicochemical features cannot fully capture;

the $\hat{\sigma}$ head learned on other proteins is therefore a less precise calibrator than on TraitGym’s shared feature pipeline. The gap is roughly $2\times$ the TraitGym Mendelian chrom-LOO marginal-bin gap (0.020, Tab. 2) and well below the cell-level worst- (k, b) gap (0.173, Tab. 10); chrom-LOO is the appropriate reference because both regimes evaluate coverage on a distribution the calibration fold did not include. (iii) *Large gap diagnostic of distribution shift*. The 0.04–0.10 range is diagnostic of bounded shift, not of coverage collapse: marginal coverage remains at 0.89 on aggregate, and the six outlier assays that drop below 0.85 have AUPRC ranging from 0.24 (CD19, base-predictor failure) to 0.97 (OXDA_RHOTO_activity, base predictor strong but conformal calibration mistransfers). HCCP degrades gracefully.

Diagnostic: how does HCCP fail on OOD outliers? We probed the abstention behaviour of HCCP on the 6 $\text{cov} < 0.85$ outlier assays vs. the 44 inliers. Three findings:

- Base-predictor quality predicts coverage but not abstention: $\text{Spearman}(\text{AUPRC}, \text{cov}) = 0.46$ ($p = 0.001$), but $\text{Spearman}(\text{AUPRC}, \text{singleton frac}) = -0.007$ ($p = 0.96$).
- Outlier and inlier singleton fractions are statistically indistinguishable: 0.378 ± 0.014 (outlier) vs. 0.355 ± 0.026 (inlier), $\Delta = 0.023$ (within one inlier std).
- Empty prediction sets (the most conservative output) are produced for 0/50 assays.

Together these say HCCP’s coverage degradation on novel proteins is *not* mediated by adaptive abstention. Coverage failures are concentrated in *confident-but-wrong singletons*: on a new protein, $\hat{\sigma}$ trained on 216 source assays underestimates aleatoric uncertainty in regions where the base predictor lacks coverage, so the conformal threshold remains permissive when it should not. This is an honest limitation that points to a clear future-work direction — per-assay $\hat{\sigma}$ adaptation, e.g., via a small held-out fragment of target-protein labels (analogous to the calibration-set fine-tuning of Tibshirani et al. [2019] for covariate shift). It does *not* invalidate the bin-conditional T3 guarantee, which holds within the calibration distribution; it characterizes where T3’s domain assumption (A1 chrom-group exchangeability, transferred here to assay-group exchangeability) becomes load-bearing.

Why this matters for the main paper. ProteinGym differs from TraitGym along three orthogonal axes: benchmark domain (protein fitness vs. variant pathogenicity), feature pipeline (physicochemical vs. pre-computed genomic foundation models), and label taxonomy (DMS fitness bins vs. ClinVar / GWAS). HCCP’s coverage degrades gracefully on both benchmarks — marginal coverage within one finite-sample slack of target on aggregate, with locally bounded $\hat{\sigma}$ -bin gap distribution (median 0.042 on ProteinGym, 0.020 marginal-bin on TraitGym Mendelian chrom-LOO) — consistent with the theoretical guarantees (T1 marginal, T2 class-cond, T3 $\hat{\sigma}$ -bin-local) being properties of the conformal calibration rather than artifacts of TraitGym alone.

E.1 Feature engineering details

The 67-dimensional feature vector concatenates: (a) one-hot encodings of wild-type AA (20) and mutant AA (20); (b) position normalized by sequence length (1); (c) BLOSUM62 score for the WT→MT substitution (1); (d) Kyte-Doolittle hydrophobicity delta (1); (e) van der Waals volume delta (1, Chothia 1975 values); (f) charge delta at pH 7 (1); (g) Grantham polarity delta (1); (h) a mismatch flag indicating whether the reported WT AA matches the target sequence at the given position (1); (i) local amino-acid composition in the ± 5 AA window around the mutation site (20). Total: $20 + 20 + 7 + 20 = 67$. Multi-mutants (e.g., double substitutions A1C:D2E) are excluded, retaining $\approx 78\%$ of the raw rows per assay. Extractor code: `scripts/37_proteingym_features.py`.

F Synthetic conditional-coverage validation

In this appendix we validate the T3 (bin-conditional) and T5 (oracle K rate) theorems on synthetic data where the oracle variance function $\sigma^*(x)$ is known by construction. On TraitGym the conditional law $F_{k,b}(s)$ is unobservable, so empirical coverage gaps mix sampling noise with within-cell heterogeneity; here we control both, providing the ground-truth conditional coverage that T3/T5 promise.

1095 F.1 Synthetic data-generating process

1096 We generate $(X, Y) \in \mathbb{R}^{d_{\text{syn}}} \times \{0, 1\}$ from

$$\begin{aligned} X &\sim \mathcal{N}(0, I_{d_{\text{syn}}}), \\ \sigma^*(x) &= \sigma_0 + \sigma_1 \cdot |x_1|, \\ Y \mid X = x &\sim \text{Bernoulli}(\text{sigmoid}(\beta^\top x + \epsilon \cdot \sigma^*(x))), \quad \epsilon \sim \mathcal{N}(0, 1), \end{aligned}$$

1097 with $d_{\text{syn}} \in \{8, 32, 128\}$, $\sigma_0 = 0.1$, $\sigma_1 = 0.4$, $\pi = P(Y = 1)$ controlled by an intercept term to
 1098 attain prevalence $\pi \in \{0.1, 0.5\}$. The first feature x_1 governs heteroscedasticity by construction, so
 1099 $\sigma^*(x)$ is a 1-D function and the Lipschitz constant L_F of the score CDF can be derived analytically
 1100 from the Gaussian-shift family.

1101 The generator and full experiment runner are implemented as `T_tools/synthetic_n_sweep.py`;
 1102 results below are from a single run of that script (full grid in < 5 CPU-min).

1103 F.2 Empirical validation of T5 (rate $O(n^{-1/2})$)

1104 We sweep $n \in \{500, 1k, 2k, 4k, 8k, 16k, 32k\}$, $\pi_{\min} \in \{0.1, 0.5\}$, $d_{\text{syn}} \in \{8, 32, 128\}$, with 5
 1105 replicates per cell. The generator code (`T_tools/synthetic_n_sweep.py`) is reproducible and
 1106 runs in under 5 CPU-min for the full grid.

Table 26: Worst-cell coverage gap $G(\hat{K}_{\text{CV}})$ (mean $\pm$ SE over 5 replicates) on synthetic data, $L_F = 4$. The dimension-free claim of T5.1 is visible as approximately matching values across d at fixed $(n, \pi_{\min})$; the $\pi_{\min} = 0.50$ rows are smaller than $\pi_{\min} = 0.10$ at matched n , with empirical compensated-constant ratio $\bar{C}_{0.10}/\bar{C}_{0.50} = 1.62$ vs the T5.1 asymptotic upper bound $\sqrt{5} \approx 2.24$ (T5.1 bounds G from above, so empirical ratios below $\sqrt{5}$ are consistent with the bound).

n	$\pi_{\min} = 0.10$			$\pi_{\min} = 0.50$		
	$d=8$	$d=32$	$d=128$	$d=8$	$d=32$	$d=128$
500	—	0.092 ± 0.011	0.094 ± 0.020	0.097 ± 0.020	0.075 ± 0.008	0.099 ± 0.016
1000	—	0.087 ± 0.009	0.088 ± 0.012	0.049 ± 0.004	0.059 ± 0.006	0.056 ± 0.005
2000	—	0.047 ± 0.006	0.057 ± 0.012	0.040 ± 0.005	0.033 ± 0.005	0.036 ± 0.005
4000	—	0.063 ± 0.012	0.038 ± 0.010	0.032 ± 0.008	0.040 ± 0.006	0.021 ± 0.003
8000	—	0.046 ± 0.008	0.018 ± 0.003	0.021 ± 0.004	0.024 ± 0.003	0.018 ± 0.002
16000	0.030 ± 0.008	0.020 ± 0.003	0.023 ± 0.007	0.019 ± 0.003	0.016 ± 0.003	0.015 ± 0.002
32000	0.023 ± 0.003	0.023 ± 0.007	0.013 ± 0.002	0.013 ± 0.002	0.010 ± 0.001	0.008 ± 0.001

1107 The log-log slope of $G(\hat{K}_{\text{CV}})$ versus n , fit per-curve across $n \geq 1,000$ to avoid the small-sample
 1108 fallback regime, is -0.44 ± 0.08 (mean $\pm$ std over the 6 fits across $d \in \{8, 32, 128\}$ and $\pi_{\min} \in$
 1109 $\{0.1, 0.5\}$); the T5.1 prediction of -0.50 lies within 1 standard deviation of the empirical mean,
 1110 consistent with the predicted rate. Figure 9 visualises the result. Figure 9(b) additionally compensates
 1111 for the rate by plotting $G(n) \cdot \sqrt{n}$, which T5.1 predicts to be approximately constant in n ; the
 1112 empirical constants are $\bar{C}_{0.10} = 3.00$ and $\bar{C}_{0.50} = 1.86$, with ratio 1.62 versus the T5.1 asymptotic
 1113 upper bound ratio $\sqrt{5} \approx 2.24$ (T5.1 is an upper bound, so empirical ratios $\leq \sqrt{5}$ are consistent).

1114 F.3 T3.b ($\hat{\sigma}$ -perturbation) ground-truth validation

1115 We inject controlled multiplicative noise into the plug-in $\hat{\sigma} = \sigma^* \cdot (1 + \eta\xi)$ for $\xi \sim \mathcal{N}(0, 1)$ and
 1116 $\eta \in \{0.05, 0.10, 0.25, 0.50\}$ at $n = 8,000$, $d = 32$, $\pi_{\min} = 0.10$, $K = 3$.

1117 The drift on class 1 (minority) is consistently larger than class 0, as expected from the smaller
 1118 $n_{\min}$ in the minority cell. The drift remains bounded under increasing η , confirming the “no cliff”
 1119 graceful-degradation property of T3.b. The empirical drifts are well within T3.b’s analytical bound
 1120 ($\sim 2\eta\bar{\sigma}/\Delta$ would be ≥ 0.5 for $\eta = 0.5$), suggesting the bound is loose but not vacuous in this
 1121 regime.




Figure 9: Synthetic n -sweep validating the $O(n^{-1/2})$ rate of Theorem 3. Worst-cell gap $G(\hat{K}_{CV})$ vs calibration size n on log-log axes, split by $\pi_{\min}$ (left 0.10, right 0.50); three coloured curves per panel cover feature dimensions $d \in \{8, 32, 128\}$ with ± 1 SE bars over 5 replicates. Y-axes are shared, so the $\pi_{\min}^{-1/2}$ vertical offset is visible at a glance: left-panel curves sit roughly $\bar{C}_{0.10}/\bar{C}_{0.50} = 1.62\times$ above the right-panel curves at matched n , below the T5.1 asymptotic upper-bound ratio $\sqrt{5} \approx 2.24$ (T5.1 is an upper bound on G , so empirical ratios below $\sqrt{5}$ are consistent with the bound). Black dashed reference: $G = \bar{C}(\pi_{\min})/\sqrt{n}$ where $\bar{C}(\pi_{\min})$ is the empirical compensated mean $G(n) \cdot \sqrt{n}$ averaged over $n \geq 1,000$ and d within each panel; this descriptive reference shows where the empirical $-1/2$ slope passes. Per-panel slope fits annotated in upper-right boxes; cross-panel summary in the title gives slope -0.44 ± 0.08 (mean $\pm$ std over 6 fits), with T5.1 prediction -0.50 within 1 standard deviation. Gray bands flag the $n < 1,000$ small-sample fallback regime excluded from the fit. The dimension-free claim is visible as overlapping d -curves within each panel.

Table 27: T3.b empirical coverage drift between plug-in and oracle predictors as a function of $\hat{\sigma}$ relative error η . Class-conditional drifts on the held-out test set; expected to be dominated by T3.b’s bound $2\eta\hat{\sigma}/\Delta + 1/(n_{\min} + 1)$. All drifts ≤ 0.022 , confirming graceful degradation; no cliff is observed even at $\eta = 0.50$.

η	$ P(\hat{\sigma}) - P(\sigma^*) $, class 0	class 1
0.05	0.0055	0.0101
0.10	0.0055	0.0213
0.25	0.0039	0.0202
0.50	0.0016	0.0180

1122 G Open Targets cross-platform replication

1123 To rule out the possibility that HCCP’s $\hat{\sigma}$ -bin gap advantage is an artifact of TraitGym’s pre-computed
 1124 pathogenicity feature pipeline (CADD + GPN-MSA + Borzoi), we replicate the headline H_2H on
 1125 a third dataset constructed from **Open Targets Platform 25.06** [Ochoa et al., 2023] fine-mapped
 1126 GWAS credible sets, using a completely independent feature pipeline (gnomAD allele frequency +
 1127 VEP consequence + AlphaMissense + SIFT + GERP + LOFTEE).

1128 **Dataset construction.** From the 2.6M credible sets in OT 25.06, we keep GWAS studies with
 1129 credible set size $\in [5, 200]$ and binarize variants by posterior inclusion probability (PIP): variants
 1130 with $\text{PIP} \geq 0.5$ are positives (24,895 unique high-PIP variants survive the AF-non-NaN filter). For
 1131 each positive, we sample 9 AF-matched random hg38 variants from the gnomAD pool *outside any*

1132 *OT credible set*, with allele frequency within ± 0.05 of the positive (TraitGym matched-9 protocol).
 1133 After deduplication on `variantId`, we subsample to 1,140 positives + 10,260 controls to mirror
 1134 TraitGym Complex matched_9 ($n = 11,400$, $\pi_+ = 0.10000$), with chrom-LOO test split chr
 1135 $\{17, 18, 19, 20, 21, 22, X\}$ matching §6. Variant overlap with TraitGym Complex / Mendelian test
 1136 sets is ≈ 0 at (chrom, pos, ref, alt) resolution and 1,846/25,068 at (chrom, pos) resolution; OT
 1137 positives are predominantly disjoint from TraitGym.

1138 **Feature pipeline.** We use 13 leak-free features extracted from OT `variant.parquet`: gnomAD
 1139 allele frequencies (af_overall, af_nfe, af_max), GERP conservation, AlphaMissense, SIFT,
 1140 FoldX, LOFTEE LoF flag, VEP-derived LoF flag, ordinal consequence severity (0–15 per the SO
 1141 hierarchy), and three feature-presence flags (NaN-handling). pip and other fine-mapping outputs are
 1142 excluded as label-leaking. The aggregator is HistGradientBoosting (max_depth=3, max_iter=300,
 1143 class_weight="balanced") trained chrom-LOO with OOF residuals for the heteroscedastic head,
 1144 identical to the TraitGym recipe. Full pool OOF AUPRC = 0.397, AUROC = 0.831 — a hard
 1145 regime comparable to TraitGym Complex (AUPRC = 0.353).

1146 **Headline H2H result.** Tab. 28 reports the four-method comparison at the recommended operating
 1147 point ($K_{\text{eval}} = 5$, $\alpha = 0.10$) under the same chrom-LOO + nested-CV $\hat{K}(c_{\text{outer}})$ protocol used for
 1148 TraitGym (§6.2, Tab. 10).

Table 28: Open Targets matched_9 H2H, $K_{\text{eval}} = 5$, $\alpha = 0.10$. HCCP partition K_{HCCP} selected per outer fold by nested chrom-LOO. Paired bootstrap on $B = 200$ chrom-resamples (seed 42); win % counts resamples on which HCCP attained a strictly smaller $\hat{\sigma}$ -bin gap than the baseline. The three baselines all hit $\hat{\sigma}$ -bin gap 0.900, the worst-cell maximum-deviation ceiling (one minority cell at near-zero coverage); the $16.7\times$ ratio is therefore better read as “HCCP holds the gap floor while the baselines saturate the ceiling” than as a multiplicative effect-size claim — the binding signal is the absolute gap of 0.054 vs. 0.900 on the worst cell, plus the 93% paired-resample win rate. Source: R_raw/cp_baselines_h2h/open_targets_*.json.

Method	Marg. cov.	Cov Y=1	$\hat{\sigma}$ -bin gap ↓	Paired Win %	p_{one}	Singleton
HCCP (ours)	0.898	0.891	0.054	—	—	0.397
RLCP	0.953	0.815	0.900	93.0%	0.075	0.611
Weighted CP	0.901	0.902	0.900	93.0%	0.075	0.507
SC-CP	0.900	0.873	0.900	93.0%	0.075	0.524

1149 HCCP attains $\hat{\sigma}$ -bin gap 0.054 vs. 0.900 for all three contemporary baselines (the worst-cell ceiling)
 1150 — a $16.7\times$ improvement at the recommended $K_{\text{eval}} = 5$; Tab. 29 reports a $12.9\text{--}16.7\times$ range across
 1151 $K_{\text{eval}} \in \{2, 5\}$ within the bias regime, of comparable magnitude to TraitGym Complex’s $13.7\times$
 1152 over the on-axis baseline (Tab. 10). The minority-coverage advantage is also preserved: HCCP and
 1153 weighted CP both hit the 0.85–0.90 band, while RLCP and SC-CP undercover the minority class.
 1154 RLCP achieves marginal coverage $0.953 > 1 - \alpha$ at the cost of $\text{cov}_{|Y=1} = 0.815$ (below the 0.85
 1155 floor) and a wide 61.1% singleton rate that exhausts HCCP’s abstention budget.

1156 **K_{eval} reversal replicates the TraitGym Mendelian boundary.** We sweep $K_{\text{eval}} \in \{2, 3, 5, 8, 10\}$
 1157 under matching nested-CV partition selection (Tab. 29). HCCP dominates by $12.9\text{--}16.7\times$ at $K_{\text{eval}} \in$
 1158 $\{2, 5\}$ (per-cell-minority ≥ 228 , deep in the bias regime), ties saturate at $K_{\text{eval}} = 8$ (per-cell-minority
 1159 ≈ 142), and *reverses* at $K_{\text{eval}} = 10$ where SC-CP overtakes HCCP (0.65 vs. 0.90, paired win rate
 1160 8.0%, $p_{\text{one}} = 0.92$ — HCCP loses with confidence). The $K_{\text{eval}} = 3$ row is anomalously high
 1161 (0.72) because nested-CV selected $K_{\text{HCCP}} = 8$ for that fold: the metric re-bins at $K_{\text{eval}} = 3$ and
 1162 aggregates over HCCP’s $K_{\text{HCCP}} = 8$ partition cells, some of which trigger the small-cell pooled
 1163 fallback. The bias-variance crossover boundary on OT sits at per-cell-minority $\approx 100\text{--}120$, consistent
 1164 with the ≈ 70 floor we report on TraitGym Mendelian (§6.2). The minority-floor argument of
 1165 Eq. (3)’s variance-term scaling is therefore consistent with the regime boundary observed across two
 1166 completely independent benchmarks; the asymptotic $K^* \approx 22$ from Eq. (3) itself is well above the
 1167 partition granularities that bind in finite samples (see also §5.2).

1168 **Cross-platform replication: three-axis takeaway.** Open Targets is a feature-pipeline-disjoint,
 1169 dataset-disjoint, and benchmark-disjoint replication of the TraitGym H2H. The headline HCCP

Table 29: K_{eval} sensitivity on Open Targets matched_9. “HCCP win ratio” is best-baseline $\hat{\sigma}$ -bin gap divided by HCCP’s; values > 1 favour HCCP. Per-cell minority = $\pi_+ n / K_{\text{eval}}$ with $\pi_+ = 0.10$, $n = 11,400$. The $K_{\text{eval}} = 3$ row reflects a $K_{\text{HCCP}} \neq K_{\text{eval}}$ partition-metric mismatch (see text); all other rows have K_{HCCP} modal $\in \{5, 8, 10\}$ aligned with K_{eval} . Excluding the $K_{\text{eval}} = 3$ mismatch row, the win ratio is monotone non-increasing in K_{eval} (equivalently, in per-cell minority count), tracking the variance-term scaling.

K_{eval}	per-cell min	HCCP gap	Best baseline gap	HCCP win ratio	Paired win %	p_{one}
2	570	0.036	0.468 (SC-CP)	12.9 $\times$	100.0%	0.005
3	380	0.718	0.900 (3-way tie)	1.25 $\times^*$	—	—
5	228	0.054	0.900 (3-way tie)	16.7 $\times$	93.0%	0.075
8	142	0.900	0.900 (4-way tie)	1.00 $\times$	—	—
10	114	0.900	0.650 (SC-CP)	0.72 $\times$	8.0%	0.92

win ratio (16.7 $\times$ at $K_{\text{eval}} = 5$) sits between TraitGym Complex (13.4 $\times$) and TraitGym Mendelian’s contested operating point; the $K_{\text{eval}} = 10$ reversal mirrors the TraitGym Mendelian $K_{\text{eval}} \geq 5$ reversal under the same per-cell minority threshold. We read this as direct empirical support for the finite-sample-floor reading of Eq. (3): the operative boundary between bias-dominated (HCCP wins) and variance-dominated (weighted CP / SC-CP wins) regimes is the per-cell quantile-estimator floor around $n_{\min} \approx 100$, not the asymptotic $\rho = 1$ crossover at $K^* \approx \sqrt{L_F R \pi_{\min} n}$ (which is ≥ 22 at TraitGym / OT scales, well above any partition we operate at in finite samples). End-to-end pipeline source: scripts/40_open_targets_download.py through scripts/44_open_targets_matched_controls.py, T_tools/open_targets_subsample.py, T_tools/cp_baselines_h2h.py -dataset open_targets.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract states the three contributions (T5.1/T5.2 finite-sample rate within the equi-bin Mondrian- K class with $\pi_{\min}^{-1/2}$ explicit; HCCP framework with T3 / T3’ guarantees; empirical validation on TraitGym + Open Targets + ProteinGym + synthetic), and §1 expands these with explicit scope qualifiers (“of the methods we evaluated”, “to our knowledge first”, “within the equi-bin Mondrian- K class”). The Mendelian $K_{\text{eval}} \geq 5$ regime where HCCP loses to weighted CP is disclosed in the abstract, as is the variant- and feature-disjoint Open Targets matched-9 replication that reproduces the 16.7 $\times$ HCCP win and the K_{eval} reversal at the predicted per-cell-minority threshold (App. G).

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: §7 “Operational envelope and limitations” enumerates four out-of-envelope failure regimes (R1 Mendelian $K_{\text{eval}} \geq 5$ at per-cell minority < 70 ; R2 Skeletal/connective OMIM cluster $\text{cov}_{|Y=1} = 0.745$, also in §6.3; R3 M $\rightarrow$ C cross-dataset joint failure with marginal coverage 0.738 within 0.004 of the Barber 2023 Thm. 2 proxy bound; R4 ProteinGym 6/50 outliers) plus four scope limitations (TV proxy approximation; large-sample K^* vs. finite-sample binding; DEGU CNN end-to-end out of scope; per-assay $\hat{\sigma}$ adaptation as future work); the K_{eval} sensitivity sweep (App. C.6) reverses the H2H story on Mendelian at $K_{\text{eval}} \geq 5$ and is reported transparently.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

1208 Answer: [\[Yes\]](#)

1209 Justification: All four theorems (T3, T3', T5.1, T5.2) state assumptions explicitly via the
1210 labeled 1/2/3/4 hierarchy. Proofs are in App. A, and the main paper §5 contains proof
1211 sketches. Auxiliary lemmas (T1, T2, T1', T2', T3-loc, T3.b, T4) and their full proofs are
1212 also in App. A.

1213 **4. Experimental result reproducibility**

1214 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
1215 perimental results of the paper to the extent that it affects the main claims and/or conclusions
1216 of the paper (regardless of whether the code and data are provided or not)?

1217 Answer: [\[Yes\]](#)

1218 Justification: Datasets (TraitGym matched_9 [Benegas et al., 2025], ProteinGym [Notin
1219 et al., 2023], Open Targets Platform 25.06 fine-mapped credible sets [Ochoa et al., 2023])
1220 are public; chrom-LOO splits ($\{17-22, X\}$) are TraitGym-canonical and reused for the
1221 Open Targets matched-9 replication; aggregator (HistGradientBoostingClassifier depth 2,
1222 100 iterations, class-balanced), $\hat{\sigma}$ head (same backbone, Gaussian NLL on chrom-LOO OOF
1223 residuals), nested chrom-LOO K -selection protocol ($K_{\text{eval}} = 3/5$ keeping ≥ 100 minority
1224 per cell), and bootstrap protocol ($B = 200$ chrom-resamples) are all specified in §4, §6, and
1225 App. B. Open Targets matched-9 dataset construction (PIP ≥ 0.5 positives, AF-matched neg-
1226 atives outside any credible set, 13-d gnomAD/VEP/AlphaMissense/SIFT/GERP/LOFTEE
1227 features, $n = 11,400$ at $\pi_+ = 0.10$) is fully documented in App. G with end-to-end pipeline
1228 scripts.

1229 **5. Open access to data and code**

1230 Question: Does the paper provide open access to the data and code, with sufficient instruc-
1231 tions to faithfully reproduce the main experimental results, as described in supplemental
1232 material?

1233 Answer: [\[Yes\]](#)

1234 Justification: TraitGym matched_9, ProteinGym, and Open Targets Platform 25.06 are
1235 publicly distributed by their creators (the Open Targets bulk parquets are mirrored on EBI
1236 FTP at <ftp.ebi.ac.uk/pub/databases/opentargets/platform/25.06/output/>).
1237 Anonymized code (aggregator training, HCCP calibration, nested-CV K -selection,
1238 bootstrap) is uploaded as a supplementary archive (supplementary.zip) accom-
1239 panying this submission, with end-to-end pipeline scripts inline-referenced from
1240 the main text: T_tools/cp_baselines_h2h.py for the H2H sweep of Tab. 15,
1241 T_tools/trait_cluster_aggregate.py for the trait-clinical clustering of Tab. 17,
1242 T_tools/synthetic_n_sweep.py for the synthetic rate validation of App. F, and
1243 scripts/40-44_open_targets_*.py + T_tools/open_targets_subsample.py for
1244 the cross-platform Open Targets replication of App. G. Permissive-license public release on
1245 acceptance.

1246 **6. Experimental setting/details**

1247 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1248 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1249 Answer: [\[Yes\]](#)

1250 Justification: §6 setup paragraph specifies splits, features, AUPRC weighting, $\alpha = 0.10$, and
1251 K -selection protocol; App. B provides the full HistGB hyperparameters (depth, iterations,
1252 class-balanced, default subsample/max_features, random_state inactive at the configuration
1253 used), $\hat{\sigma}$ -bin edge construction (calibration-fold quantiles), and the small-cell fallback
1254 ($n_{k,b} < \lceil 1/\alpha \rceil$ triggers pooled T2 threshold).

1255 **7. Experiment statistical significance**

1256 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1257 information about the statistical significance of the experiments?

1258 Answer: [\[Yes\]](#)

1259 Justification: Tab. 2 HCCP rows report $\pm$ one std from $B = 200$ chromosome-level bootstrap
1260 resamples; Tab. 9 and Tab. 10 report per-method bootstrap mean $\pm$ std. Bootstrap captures

1261 calibration-set variability (chrom-resample with replacement). Single-seed point estimates
1262 are flagged as such in Tab. 15 (Mendelian $K_{\text{eval}} = 3$ HCCP gap 0.163 point estimate vs
1263 0.173 ± 0.192 bootstrap mean in Tab. 10) and in Tab. 22; the offset between the two is
1264 documented in the caption of Tab. 15.

1265 8. Experiments compute resources

1266 Question: For each experiment, does the paper provide sufficient information on the com-
1267 puter resources (type of compute workers, memory, time of execution) needed to reproduce
1268 the experiments?

1269 Answer: [Yes]

1270 Justification: App. B (compute table) reports: HCCP pipeline (aggregator + $\hat{\sigma}$ head +
1271 Mondrian calibration) runs entirely on CPU 16-core; total ≤ 4 CPU-hour per feature set
1272 per dataset; $B = 200$ chrom-bootstrap adds ~ 30 CPU-min total. ProteinGym 50-assay
1273 LOO runs on CPU 16-core in < 10 CPU-hour. The Open Targets cross-platform replication
1274 (App. G) downloads ≈ 6 GB from EBI FTP (one-time ~ 30 min on a typical link), parses
1275 to $\approx 13\text{M}$ variant rows in ≈ 5 CPU-min, builds the matched-9 subset in ≈ 4 CPU-min,
1276 trains chrom-LOO scores in ≈ 30 CPU-sec, and runs the H2H + paired bootstrap in ≈ 30
1277 CPU-sec at $K_{\text{eval}} = 5$, $B = 200$. No GPU is required for the headline pipeline.

1278 9. Code of ethics

1279 Question: Does the research conducted in the paper conform, in every respect, with the
1280 NeurIPS Code of Ethics?

1281 Answer: [Yes]

1282 Justification: We have read the NeurIPS Code of Ethics. The paper uses publicly available
1283 variant-effect benchmark datasets (TraitGym, ProteinGym, Open Targets Platform) under
1284 their respective open licences; no human-subject data collection or crowdsourcing was
1285 performed. Clinical-deployment risks of variant prediction are discussed in §7.

1286 10. Broader impacts

1287 Question: Does the paper discuss both potential positive societal impacts and negative
1288 societal impacts of the work performed?

1289 Answer: [Yes]

1290 Justification: §7 broader-impact paragraph discusses (i) clinical deployment risks of variant
1291 prediction and the need to integrate HCCP singletons into ACMG/AMP multi-evidence
1292 frameworks rather than as stand-alone diagnostics; (ii) population representativeness limita-
1293 tions of TraitGym (over-represents European-ancestry individuals from ClinVar / GWAS
1294 catalog), shared by Open Targets which is also harmonized from European-majority cohorts
1295 (UK Biobank, FinnGen, GWAS catalog); (iii) ancestry-stratified Mondrian as a future-work
1296 direction for equitable deployment.

1297 11. Safeguards

1298 Question: Does the paper describe safeguards that have been put in place for responsible
1299 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1300 image generators, or scraped datasets)?

1301 Answer: [N/A]

1302 Justification: HCCP is a post-hoc conformal-calibration framework consuming an arbitrary
1303 classifier $\hat{p}$ and variance head $\hat{\sigma}$. It is not a pre-trained generative model, language model,
1304 or scraped dataset, and we do not release any high-risk asset. All datasets used (TraitGym,
1305 ProteinGym, Open Targets Platform) are already public and were created by their original
1306 authors with their own release decisions.

1307 12. Licenses for existing assets

1308 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1309 the paper, properly credited and are the license and terms of use explicitly mentioned and
1310 properly respected?

1311 Answer: [Yes]

1312 Justification: TraitGym [Benegas et al., 2025], ProteinGym [Notin et al., 2023], Open Targets
1313 Platform 25.06 [Ochoa et al., 2023], CADD [Kircher et al., 2014], Borzoi [Linder et al.,
1314 2025], GPN-MSA [Benegas et al., 2023], AlphaGenome [Avsec et al., 2026] are all cited
1315 via the bibliography with full publication metadata; the crepes package [Boström, 2022]
1316 is used for the H2H baselines and acknowledged. We use each asset in accordance with
1317 its publicly distributed licence (TraitGym: MIT; ProteinGym: CC-BY-4.0; Open Targets
1318 Platform parquets: CC0-1.0 / open-data licence per the EBI distribution; CADD: free
1319 academic; Borzoi/GPN-MSA: open-source per the original repos; AlphaGenome: Google
1320 API terms; crepes: BSD-3).

1321 13. New assets

1322 Question: Are new assets introduced in the paper well documented and is the documentation
1323 provided alongside the assets?

1324 Answer: [Yes]

1325 Justification: New assets (HCCP calibration code, A-SL audit utility, bootstrap CP baseline
1326 runner) are anonymized and supplied as supplementary material; App. B documents the
1327 script entry-points, expected inputs / outputs, and runtime estimates. Final release on
1328 acceptance will follow standard open-source documentation templates.

1329 14. Crowdsourcing and research with human subjects

1330 Question: For crowdsourcing experiments and research with human subjects, does the paper
1331 include the full text of instructions given to participants and screenshots, if applicable, as
1332 well as details about compensation (if any)?

1333 Answer: [N/A]

1334 Justification: The paper does not involve crowdsourcing or research with human subjects.
1335 All datasets used are publicly distributed variant-effect benchmarks (TraitGym derived from
1336 ClinVar / GWAS catalog; ProteinGym from published deep-mutational-scanning assays;
1337 Open Targets Platform from harmonized GWAS catalog / FinnGen / UK Biobank Pharma
1338 Proteomics summary statistics with SuSiE / PICS fine-mapping).

1339 15. Institutional review board (IRB) approvals or equivalent for research with human 1340 subjects

1341 Question: Does the paper describe potential risks incurred by study participants, whether
1342 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1343 approvals (or an equivalent approval/review based on the requirements of your country or
1344 institution) were obtained?

1345 Answer: [N/A]

1346 Justification: The paper does not conduct human-subjects research. Datasets are public
1347 benchmarks created and released by their original authors under ethics frameworks of the
1348 originating institutions; we use the released aggregate variant labels and do not access any
1349 individual-level data.

1350 16. Declaration of LLM usage

1351 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1352 non-standard component of the core methods in this research? Note that if the LLM is used
1353 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1354 scientific rigor, or originality of the research, declaration is not required.

1355 Answer: [N/A]

1356 Justification: LLMs were not used as any component of the core methodology. The aggrega-
1357 tor $\hat{p}$ is a gradient-booster classifier and the variance head $\hat{\sigma}$ is a gradient-booster regressor;
1358 both consume pre-computed numeric features (CADD, Borzoi, GPN-MSA scores) rather
1359 than language.